COMP0086 Summative Assignment

Nov 14, 2022

Question 1

(a) Our sample space for images is $\{0,1\}^D$, a binary space with D dimensions, the number of pixels in the image. Thus, picking the exponential family best suited on this sample space is the D-dimensional multivariate Bernoulli distribution which shares the same sample space. On the other hand, a D-dimensional multivariate Gaussian has the sample space $\mathbb{R}^D$, which does not match the sample space of our data. To match our data sample space, we would have to define additional mapping between our data and model spaces which adds unnecessary complexity. Thus it would be inappropriate to model this dataset of images with a multivariate Gaussian.

(b) For $\mathcal{D} := \{x^{(n)}\}_{n=1}^N$ a data set of N images, the joint likelihood (assuming images are independently and identically distributed) is the product of N D-dimensional multivariate Bernoulli distributions, one for each image:

$$P(\mathcal{D}|\mathbf{p}) = \prod_{n=1}^N P(x^{(n)}|\mathbf{p})$$

Substituting the D-dimensional multivariate Bernoulli:

$$P(\mathcal{D}|\mathbf{p}) = \prod_{n=1}^N \prod_{d=1}^D p_d^{x_d^{(n)}} (1 - p_d)^{1 - x_d^{(n)}}$$

Taking the logarithm of this, we get the log likelihood:

$$\mathcal{L}(\mathcal{D}|\mathbf{p}) = \sum_{n=1}^N \sum_{d=1}^D [x_d^{(n)} \log(p_d) + (1 - x_d^{(n)}) \log(1 - p_d)]$$

Note that since the logarithm of the mean is a monotonic increasing function on $\mathbb{R}_+$, the maximisers and minimisers of the likelihood do not change.

To solve for the maximum likelihood estimate, $\hat{p}_d$, we can take the derivative of $\mathcal{L}(\mathcal{D}|\mathbf{p})$ with respect to p_d, the d^{th} element of $\mathbf{p}$:

$$\frac{\partial \mathcal{L}(\mathcal{D}|\mathbf{p})}{\partial p_d} = \sum_{n=1}^N \left(\frac{x_d^{(n)}}{p_d} - \frac{1 - x_d^{(n)}}{1 - p_d} \right)$$

1

$$\frac{\partial \mathcal{L}(\mathcal{D}|\mathbf{p})}{\partial p_d} = \frac{\sum_{n=1}^{N} x_d^{(n)}}{p_d} - \frac{\sum_{n=1}^{N}(1 - x_d^{(n)})}{1 - p_d}$$

and set the derivative to zero and solve for $\hat{p}_d$:

$$\frac{\sum_{n=1}^{N} x_d^{(n)}}{\hat{p}_d} - \frac{\sum_{n=1}^{N}(1 - x_d^{(n)})}{1 - \hat{p}_d} = 0$$

$$\sum_{n=1}^{N} x_d^{(n)} - \hat{p}_d \sum_{n=1}^{N} x_d^{(n)} - \hat{p}_d \cdot N + \hat{p}_d \sum_{n=1}^{N} x_d^{(n)} = 0$$

$$\hat{p}_d = \frac{1}{N} \sum_{n=1}^{N} x_d^{(n)}$$

Moreover, the second derivative with respect to p_d:

$$\frac{\partial \mathcal{L}(\mathcal{D}|\mathbf{p})}{\partial p_d^2} = \frac{-\sum_{n=1}^{N} x_d^{(n)}}{p_d^2} + \frac{-\sum_{n=1}^{N}(1 - x_d^{(n)})}{(1 - p_d)^2}$$

For a maximum, we need to show that the second derivative is negative. Since $x_d^{(n)} \in \{0, 1\}$, in the worst case, of $N = 1$, the single pixel must either be white ($\sum_{n=1}^{N} > 0$) or black ($\sum_{n=1}^{N} 1 - x_d^{(n)} > 0$) so $\frac{\partial \mathcal{L}(\mathcal{D}|\mathbf{p})}{\partial p_d^2} < 0$ will be guaranteed and $\hat{p}_d$ is a maximum as required for the maximum likelihood estimate.

Because we assume that each pixel is independent (we are taking the product of D one dimensional Bernoulli distributions), we can express the maximum likelihood for $\mathbf{p}$ in matrix form as $\hat{\mathbf{p}}$:

$$\hat{\mathbf{p}} = \frac{1}{N} \sum_{n=1}^{N} \mathbf{x}^{(n)}$$

(c) From Bayes' Theorem:

$$P(\mathbf{p}|\mathcal{D}) = \frac{P(\mathcal{D}|\mathbf{p})P(\mathbf{p})}{P(\mathcal{D})}$$

Taking the logarithm:

$$\mathcal{L}(\mathbf{p}|\mathcal{D}) = \mathcal{L}(\mathcal{D}|\mathbf{p}) + \mathcal{L}(\mathbf{p}) - \mathcal{L}(\mathcal{D})$$

Taking the derivative with respect to p_d:

$$\frac{\partial \mathcal{L}(\mathbf{p}|\mathcal{D})}{\partial p_d} = \frac{\partial \mathcal{L}(\mathcal{D}|\mathbf{p})}{\partial p_d} + \frac{\partial \mathcal{L}(\mathbf{p})}{\partial p_d}$$

where $\frac{\partial \mathcal{L}(\mathcal{D})}{\partial p_d} = 0$ because it doesn't depend on p_d.

We know (b):

$$\frac{\partial \mathcal{L}(\mathcal{D}|\mathbf{p})}{\partial p_d} = \frac{\sum_{n=1}^{N} x_d^{(n)}}{p_d} - \frac{\sum_{n=1}^{N}(1 - x_d^{(n)})}{1 - p_d}$$

For the second term $\frac{\partial \mathcal{L}(\mathbf{p})}{\partial p_d}$, we start with $P(\mathbf{p})$):

$$P(\mathbf{p})) = \prod_{d=1}^{D} P(p_d)$$

Assuming independent Beta priors on the parameters p_d:

$$P(\mathbf{p}) = \prod_{d=1}^{D} \frac{1}{B(\alpha, \beta)} p_d^{\alpha-1}(1 - p_d)^{\beta-1}$$

Taking the logarithm:

$$\mathcal{L}(\mathbf{p}) = \sum_{d=1}^{D} -\log(B(\alpha, \beta)) + (\alpha - 1)\log p_d + (\beta - 1)\log(1 - p_d)$$

Taking the derivative with respect to p_d:

$$\frac{\partial \mathcal{L}(\mathbf{p})}{\partial p_d} = -\log(B(\alpha, \beta)) + (\alpha - 1)\log p_d + (\beta - 1)\log(1 - p_d))$$

Since we are only concerned with p_d, we are only left with a single element of the summation pertaining to p_d.

Simplifying:

$$\frac{\partial \mathcal{L}(\mathbf{p})}{\partial p_d} = \frac{(\alpha - 1)}{p_d} - \frac{(\beta - 1)}{1 - p_d}$$

Combining to have an expression for the log posterior derivative $\frac{\partial \mathcal{L}(\mathbf{p}|\mathcal{D})}{\partial p_d}$:

$$\frac{\partial \mathcal{L}(\mathbf{p}|\mathcal{D})}{\partial p_d} = \frac{\sum_{n=1}^{N} x_d^{(n)}}{p_d} - \frac{\sum_{n=1}^{N}(1 - x_d^{(n)})}{1 - p_d} + \frac{(\alpha - 1)}{p_d} - \frac{(\beta - 1)}{1 - p_d}$$

$$\frac{\partial \mathcal{L}(\mathbf{p}|\mathcal{D})}{\partial p_d} = \frac{(\alpha - 1) + \sum_{n=1}^{N} x_d^{(n)}}{p_d} - \frac{(\beta - 1) + \sum_{n=1}^{N}(1 - x_d^{(n)})}{1 - p_d}$$

To find the maximum a posteriori (MAP) estimate $\hat{p}_d$ with $\frac{\partial \mathcal{L}(\mathbf{p}|\mathcal{D})}{\partial p_d} = 0$:

$$0 = \frac{(\alpha - 1) + \sum_{n=1}^{N} x_d^{(n)}}{\hat{p}_d} - \frac{(\beta - 1) + \sum_{n=1}^{N}(1 - x_d^{(n)})}{1 - \hat{p}_d}$$

$$0 = (1 - \hat{p}_d)(\alpha - 1) + (1 - \hat{p}_d)\left(\sum_{n=1}^{N} x_d^{(n)}\right) - \hat{p}_d(\beta - 1) - \hat{p}_d\left(\sum_{n=1}^{N}(1 - x_d^{(n)})\right)$$

$$0 = (\alpha - \alpha\hat{p}_d + \hat{p}_d - 1) + \left(\sum_{n=1}^{N} x_d^{(n)} - \hat{p}_d \sum_{n=1}^{N} x_d^{(n)}\right) - (\hat{p}_d\beta - \hat{p}_d) - \left(\hat{p}_d \cdot N - \hat{p}_d \sum_{n=1}^{N} x_d^{(n)}\right)$$

Cancelling the $\hat{p}_d \sum_{n=1}^{N} x_d^{(n)}$ terms:

$$0 = \alpha - \alpha\hat{p}_d + \hat{p}_d - 1 + \sum_{n=1}^{N} x_d^{(n)} - \hat{p}_d\beta + \hat{p}_d - \hat{p}_d \cdot N$$

$$0 = \hat{p}_d(2 - \alpha - \beta - N) + \alpha - 1 + \sum_{n=1}^{N} x_d^{(n)}$$

$$\hat{p}_d = \frac{\alpha - 1 + \sum_{n=1}^{N} x_d^{(n)}}{(N + \alpha + \beta - 2)}$$

To show that this is a maximum, the second derivative is:

$$\frac{\partial^2 \mathcal{L}(\mathbf{p}|\mathcal{D})}{(\partial p_d)^2} = \frac{(1 - \alpha) - \sum_{n=1}^{N} x_d^{(n)}}{(p_d)^2} + \frac{(1 - \beta) - \sum_{n=1}^{N}(1 - x_d^{(n)})}{(1 - p_d)^2}$$

.

For a maximum, we need $\frac{\partial^2 \mathcal{L}(\mathbf{p}|\mathcal{D})}{(\partial p_d)^2} < 0$ meaning that we need at least one of the strict inequalities $\alpha < 1 - \sum_{n=1}^{N} x_d^{(n)}$ or $\beta < 1 - \sum_{n=1}^{N}(1 - x_d^{(n)})$ to be satisified, where the other can be $\leq$. The Beta distribution requires $\alpha > 0$ and $\beta > 0$ so this requirement will always be satisfied (in the worst case of a single image, either $x_d^{(1)} = 1$ or $1 - x_d^{(1)} = 1$).

Due to independence of our likelihood and priors for each dimension, we can express the maximum a priori for $\mathbf{p}$ in matrix form as $\hat{\mathbf{p}}$:

$$\hat{\mathbf{p}} = \frac{\alpha - 1 + \sum_{n=1}^{N} \mathbf{x}^n}{(N + \alpha + \beta - 2)}$$

(d&e) The Python code for MLE and MAP:

```python
import matplotlib.pyplot as plt
import numpy as np


def _compute_maximum_likelihood_estimate(x: np.ndarray) -> np.ndarray:
    """
    X: numpy array of shape (N, D)
    """
    return np.mean(x, axis=0)


def _compute_maximum_a_priori_estimate(
    x: np.ndarray, alpha: float, beta: float
) -> np.ndarray:
    """
    X: numpy array of shape (N, D)
    alpha: param of prior distribution
    beta: param of prior distribution
    """

    n, _ = x.shape
    return (alpha - 1 + np.sum(x, axis=0)) / (n + alpha + beta - 2)


def d(x, figure_path, figure_title):
    maximum_likelihood = _compute_maximum_likelihood_estimate(x)
    plt.figure()
    plt.imshow(
        np.reshape(maximum_likelihood, (8, 8)),
        interpolation="None",
    )
    plt.colorbar()
    plt.axis("off")
    plt.title(figure_title)
    plt.savefig(figure_path)


def e(x, alpha, beta, figure_path, figure_title):
    maximum_a_priori = _compute_maximum_a_priori_estimate(x, alpha, beta)
    plt.figure()
    plt.imshow(
        np.reshape(maximum_a_priori, (8, 8)),
        interpolation="None",
    )
    plt.colorbar()
    plt.axis("off")
    plt.title(figure_title)
    plt.savefig(f"{figure_path}.png")

    maximum_likelihood = _compute_maximum_likelihood_estimate(x)
    plt.figure()
    plt.imshow(
        np.reshape(maximum_a_priori - maximum_likelihood, (8, 8)),
        interpolation="None",
    )
    plt.colorbar()
    plt.axis("off")
    plt.title(f"MAP vs MLE")
    plt.savefig(f"{figure_path}-mle-vs-map.png")
```

src/solutions/q1.py

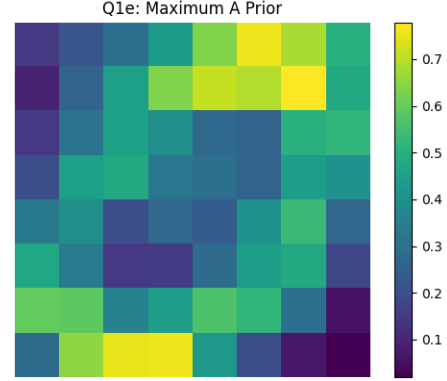Displaying the learned parameters:

Figure 1: ML parameters

Figure 2: MAP parameters

Comparing the equations:

$$\hat{\mathbf{p}}^{MLE} = \frac{1}{N} \sum_{n=1}^{N} \mathbf{x}^{(n)}$$

and

$$\hat{\mathbf{p}}^{MAP} = \frac{\alpha - 1 + \sum_{n=1}^{N} \mathbf{x}^n}{(N + \alpha + \beta - 2)}$$

We see that when the number of data points increases, the MAP estimate approaches $\frac{1}{N} \sum_{n=1}^{N} \mathbf{x}^{(n)}$, the MLE. This result makes sense because as our data set gets bigger, are less reliant on our prior. Moreover, if in our data set of images, a specific pixel in all of these images are white or all black, the MLE for that pixel will be binary. This may not be be representative of our intuitions about images, where there should be some non-zero probability of a pixel being black or white. Thus, introducing an appropriate prior will ensure that the probability of that pixel will never be zero or one. Moreover, a prior can also help ensure numerical stability during calculations. The logarithm of zero is negative infinity, so having probability zero can be problematic for log-likelihoods calculations. A prior can ensure non-zero probabilities. In our case, by using a small Beta distribution prior on each pixel, our parameter values are biased be close to 0.5 and to never be at the extremities 0 and 1. We can see this in our MAP parameters figure where the range of our parameters is smaller than the range for the ML parameters. Also, the MLE has values of zero whereas the MAP does not. Interestingly, when $\alpha = 1$ and $\beta = 1$, $\hat{\mathbf{p}}^{MLE} = \hat{\mathbf{p}}^{MAP}$, intuitively this is when the prior is a uniform distribution and so there aren't any biases on the location of $\mathbf{p}$ and we recover the MLE. On the other hand, a misspecified prior can be problematic, as the estimated parameters might be skewed by the prior and not properly represent the underlying data generating process and this can be worse than using the MLE.

We can also visualise $\hat{\mathbf{p}}^{MAP} - \hat{\mathbf{p}}^{MLE}$ to see that for likelihoods greater than 0.5 in the MLE, the MAP has a lower value and for likelihoods less than 0.5, the MAP has a higher value, confirming our intuitions.
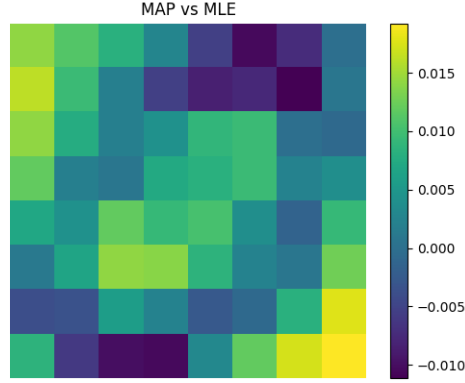
Figure 3: $\hat{\mathbf{p}}^{MAP} - \hat{\mathbf{p}}^{MLE}$

Question 2

(a) When all D components are generated from a Bernoulli distribution with $p_d = 0.5$, we have the likelihood function for model M_1:

$$P(\mathbf{x}^{(n)}|\mathbf{P}^{(1)} = [0.5, 0.5, ..., 0.5]^T) = \prod_{d=1}^{D}(0.5)^{x_d^{(n)}}(0.5)^{1-x_d^{(n)}}$$

(b) When all D components are generated from Bernoulli distributions with unknown, but identical, p_d, we have the likelihood function for model M_2:

$$P(\mathbf{x}^{(n)}|\mathbf{p}^{(2)} = [p_d, p_d, ..., p_d]^T) = \prod_{d'=1}^{D} p_{d'}^{x_{d'}^{(n)}}(1 - p_d)^{1-x_{d'}^{(n)}}$$

(c) When each component is Bernoulli distributed with separate, unknown p_d, we have the likelihood function for model M_3:

$$P(\mathbf{x}^{(n)}|\mathbf{p}^{(3)} = [p_1, p_2, ..., p_D]^T) = \prod_{d=1}^{D} p_d^{x_d^{(n)}}(1 - p_d)^{1-x_d^{(n)}}$$

For each model M_i, we can marginalise out P to get $P(\{\mathbf{x}^{(n)}\}_{n=1}^{N}|M_i)$:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^{N}|M_i) = \int_0^1 ... \int_0^1 P(\{\mathbf{x}^{(n)}\}_{n=1}^{N}|p_d, M_i)P(p_d|M_i)dp_1...dp_D$$

where $d = 1, ..., D$ and $\{\mathbf{x}^{(n)}\}_{n=1}^{N}$ is our data set.

Given that the prior of any unknown probabilities is uniform, i.e. $P(p_d|M_i) = 1$. We can simplify:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^{N}|M_i) = \int_0^1 ... \int_0^1 P(\{\mathbf{x}^{(n)}\}_{n=1}^{N}|p_d, M_i)dp_1...dp_D$$

For M_1, we have that all pixels have probability 0.5:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_1) = \int_0^1 \dots \int_0^1 \prod_{j=n}^N \prod_{d=1}^D (0.5)^{x_d^{(n)}} (1 - 0.5)^{1 - x_d^{(n)}} d\theta_1 \dots d\theta_D$$

We can remove the integrals and knowing that either $x_d^{(n)}$ or $1 - x_d^{(n)}$ will be 1 and the other zero, we can simplify $(0.5)^{x_d^{(n)}} (1 - 0.5)^{1 - x_d^{(n)}}$ to 0.5:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_1) = \prod_{j=1}^N \prod_{d=1}^D (0.5)$$

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_1) = (0.5)^{N \cdot D}$$

For M_2, we have that all pixels share some probability p_d so we only need to integrate over a single variable p_d:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_2) = \int_0^1 \prod_{n=1}^N \prod_{d'=1}^D p^{x_{d'}^{(n)}} (1 - p_d)^{1 - x_{d'}^{(n)}} dp_d$$

Changing the products to sums:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_2) = \int_0^1 p_d^{\sum_{n=1}^N \sum_{d'=1}^D x_{d'}^{(n)}} (1 - p)^{\sum_{j=1}^N \sum_{d'=1}^D 1 - x_{d'}^{(n)}} dp_d$$

Rewriting:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_2) = \int_0^1 p_d^k (1 - p_{d'=1})^{N \cdot D - k} dp_d$$

where $k = \sum_{n=1}^N \sum_{d'=1}^D x_{d'}^{(n)}$.

This integral is the beta function:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_2) = \frac{k!(N \cdot D - k)!}{(N \cdot D + 1)!}$$

For M_3, we need an integral for each p_d:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_3) = \int_0^1 \dots \int_0^1 \prod_{n=1}^N \prod_{d=1}^D p_d^{x_d^{(n)}} (1 - p_d)^{1 - x_d^{(n)}} dp_1 \dots dp_D$$

We can separate the integrals to only contain the relevant p_d:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_3) = \prod_{d=1}^D \left(\int_0^1 \prod_{n=1}^N p_d^{x_d^{(n)}} (1 - p_d)^{1 - x_d^{(n)}} dp_d \right)$$

Changing the products to sums:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N|M_3) = \prod_{d=1}^D \left(\int_0^1 p_d^{\sum_{n=1}^N x_d^{(n)}} (1-p_d)^{\sum_{n=1}^N 1-x_d^{(j)}} dp_d \right)$$

In this case, we have the product of integrals where each evaluates to a beta function:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N|M_3) = \prod_{d=1}^D \frac{k_d!(N-k_d)!}{(N+1)!}$$

where $k_d = \sum_{n=1}^N x_d^{(in}$.

The posterior probability of a model M_i can be expressed:

$$P(M_i|\{\mathbf{x}^{(n)}\}_{n=1}^N) = \frac{P(\{\mathbf{x}^{(n)}\}_{n=1}^N|M_i)P(M_i)}{P(\{\mathbf{x}^{(n)}\}_{n=1}^N)}$$

We only ave three models, so in this case the normalisation $P(D)$ can be expressed as a sum:

$$P(M_i|\{\mathbf{x}^{(n)}\}_{n=1}^N) = \frac{P(\{\mathbf{x}^{(n)}\}_{n=1}^N|M_i)P(M_i)}{\sum_{i\in\{1,2,3\}} P(\{\mathbf{x}^{(n)}\}_{n=1}^N|M_i)P(M_i)}$$

Given that $P(M_i) = \frac{1}{3}$ for all $i \in \{1,2,3\}$:

$$P(M_i|\{\mathbf{x}^{(n)}\}_{n=1}^N) = \frac{P(\{\mathbf{x}^{(n)}\}_{n=1}^N|M_i)}{\sum_{i\in\{1,2,3\}} P(\{\mathbf{x}^{(n)}\}_{n=1}^N|M_i)}$$

Calculating the posterior probabilities of each of the three models having generated the data in binarydigits.txt using python we can show the values in the table below:

i	$P(M_i|\mathcal{D})$
1	1E-1924
2	1E-1858
3	1-(1E-1924)-(1E-1858)

The Python code for calculating the posterior probabilities of the three models:

```python
import numpy as np
import pandas as pd
from scipy.special import betaln, logsumexp


def _log_p_d_given_m1(x):
    n, d = x.shape
    return n * d * np.log(0.5)


def _log_p_d_given_m2(x):
    n, d = x.shape
    k = np.sum(x, axis=0).astype(int)
    return betaln(np.sum(k) + 1, n * d - np.sum(k) + 1)


def _log_p_d_given_m3(x):
    n, _ = x.shape
    k = np.sum(x, axis=0).astype(int)
    return logsumexp(betaln(k + 1, n - k + 1))


def c(x, table_path):
    log_p_d_given_m = np.array(
        [
            _log_p_d_given_m1(x),
            _log_p_d_given_m2(x),
            _log_p_d_given_m3(x),
        ]
    )
    log_p_m_given_d = log_p_d_given_m - logsumexp(log_p_d_given_m)
    df = pd.DataFrame(
        data=np.array(
            [
                np.arange(len(log_p_m_given_d)).astype(int) + 1,
                [f"1E{int(x/np.log(10))}" for x in log_p_m_given_d[:-1]]
                + [
                    f"1-{'-'.join([f'(1E{int(x/np.log(10))})' for x in log_p_m_given_d[:-1]])}"
                ],
            ]
        ).T,
        columns=["Model", "P(M_i|D)"],
    )
    df.set_index("Model", inplace=True)
    df.to_csv(table_path)
```

src/solutions/q2.py

Question 3

(a) The likelihood for a model consisting of a mixture of K multivariate Bernoulli distributions can be expressed as the product across N data points:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^{N}|\theta) = \prod_{i=1}^{N} P(x_i|\theta)$$

where $\{\mathbf{x}^{(n)}\}_{n=1}^{N}$ is our data set with $\mathbf{x}^{(n)} \in \mathbb{R}^{D\times 1}$ and $\theta = \{\pi, \mathbf{P}\}$, $\pi = [\pi_1, ..., \pi_K] \in \mathbb{R}^{K\times 1}$ our mixing proportions $(0 \leq \pi_k \leq 1; \sum_k \pi_k = 1)$ and $\mathbf{P} \in \mathbb{R}^{D\times K}$ the K Bernoulli parameter vectors with elements p_{kd} denoting the probability that pixel d takes value 1 under mixture component k. We are also the images are iid under the model, and that the pixels are independent of each other within each component distribution.

For each $P(\mathbf{x}^{(n)}|\theta)$:

$$P(\mathbf{x}^{(n)}|\theta) = \sum_{k=1}^{K} \pi_k \prod_{d=1}^{D}(p_{kd})^{\mathbf{x}_d^{(n)}}(1 - p_{kd})^{1-\mathbf{x}_d^{(n)}}$$

The log-likelihood $\mathcal{L}(\mathbf{x}^{(n)}|\theta)$ can be expressed in matrix form:

$$\mathcal{L}(\mathbf{x}^{(n)}|\theta) = \log \sum_{k=1}^{K} \pi_k \exp\left(\mathbf{x}^{(n)}\log(\mathbf{P}_k) + (1 - \mathbf{x}^{(n)})\log(1 - \mathbf{P}_k)\right)$$

which can be further vectorised using Python scipy's *logsumexp* operation.

Moreover, the log-likelihood $\mathcal{L}(\{\mathbf{x}^{(n)}\}_{n=1}^{N}|\theta)$ can be expressed:

$$\mathcal{L}(\{\mathbf{x}^{(n)}\}_{n=1}^{N}|\theta) = \sum_{i=1}^{N}\left(\log \sum_{k=1}^{K} \pi_k \exp\left(\mathbf{x}^{(n)}\log(\mathbf{P}_k) + (1 - \mathbf{x}^{(n)})\log(1 - \mathbf{P}_k)\right)\right)$$

(b) The expression for the responsibility of mixture component k for data vector $x^{(n)}$, i.e. $r_{nk} \equiv P(s^{(n)} = k|x^{(n)}, \pi, \mathbf{P})$. This computation provides the E-step for an EM algorithm.

We know that:

$$P(A|B) \propto P(B|A)P(A)$$

Thus,

$$P(s^{(n)} = k|x^{(n)}, \pi, \mathbf{P}) \propto P(x^{(n)}|s^{(n)} = k, \pi, \mathbf{P})P(s^{(n)} = k|\pi, \mathbf{P})$$

where $s^{(n)} \in \{1, ..., K\}$ a discrete latent variable where $P(s^{(n)} = k|x^{(n)}|\pi) = \pi_k$. Note that $P(s^{(n)} = k|x^{(n)}|\pi) = P(s^{(n)} = k|x^{(n)}|\pi, \mathbf{P})$ as $s^{(n)}$ isn't dependent on $\mathbf{P}$.

Let $P(s^{(n)} = k|x^{(n)}, \pi, \mathbf{P}) \propto P(s^{(n)}$ be the unnormalised responsibility $\tilde{r}_{nk}$. Using the mixture for component k and the likelihood function of component k:

$$\tilde{r}_{nk} = \pi_k \prod_{d=1}^{D} (p_{kd})^{x_d^{(n)}} (1 - p_{kd})^{1 - x_d^{(n)}}$$

Normalising across the components:

$$r_{nk} = \frac{\tilde{r}_{nk}}{\sum_{j=1}^{K} \tilde{r}_{nj}}$$

and r_{nk}, we have calculated $P(s^{(n)} = k | x^{(n)}, \pi, \mathbf{P})$ for the E step.
Moreover,

$$\log \tilde{r}_{nk} = \log \pi_k + \sum_{d=1}^{D} \left(x_d^{(n)} \log(p_{kd}) + (1 - x_d^{(n)}) \log(1 - \exp(\log(p_{kd}))) \right)$$

and

$$\log r_{nk} = \log \tilde{r}_{nk} - \log \sum_{j=1}^{K} \exp(\log \tilde{r}_{nj})$$

which can be vectorised using Python scipy's *logsumexp* operation.

(c) We know that the expectation log joint can be expressed:

$$\left\langle \sum_n \log P(x^{(n)}, s^{(n)} | \pi, \mathbf{P}) \right\rangle_{q(\{s^{(n)}\})} = \sum_{n=1}^{N} q(s^{(n)}) \log P(x^{(n)}, s^{(n)} | \pi, \mathbf{P})$$

Let this quantity be E.

Filling in the appropriate quantities:

$$E = \sum_n r_n^T [\log(\pi) + \log(\mathbf{P})^T x^{(n)} + \log(1 - \mathbf{P})^T (1 - x^{(n)})]$$

where $q(s^{(n)}) = r_n^T$ and $\log P(x^{(n)}, s^{(n)}) = \log(\pi) + \log(\mathbf{P})^T x^{(n)} + \log(1 - \mathbf{P})^T (1 - x^{(n)})$. Noting that $P(x^{(n)}, s^{(n)} | \pi, \mathbf{P}) = P(x^{(n)} | s^{(n)}) \pi, \mathbf{P}) P(s^{(n)}) | \pi, \mathbf{P}) = \tilde{r}_n$

To maximise with respect to π and P for the M step, we want to take the derivative, set to zero, and solve for $\hat{\pi}$ and $\hat{P}$.

For the k^{th} element of π:

$$\frac{\partial E}{\partial \pi_k} = \sum_n r_{nk} \frac{1}{\pi_m}$$

and

$$\frac{\partial E}{\partial \pi_k} + \lambda = 0$$

where λ is a Lagrange multiplier ensuring that the mixing proportions sum to unity. Thus,

$$\hat{\pi}_m = \frac{1}{N} \sum_n r_{nk}$$

For the dk^{th} element of P:

$$\frac{\partial E}{\partial P_{dk}} = \sum_n r_{nk} \frac{\partial}{\partial P_{dk}} [x_d^{(n)} \log P_{dk} + (1 - x_d^{(n)}) \log(1 - P_{dk})]$$

Simplying:

$$\frac{\partial E}{\partial P_{dk}} = \sum_n r_{nk} (\frac{x_d^{(n)}}{P_{dk}} - \frac{1 - x_d^{(n)}}{1 - P_{dk}})$$

Setting the derivative to zero:

$$\frac{\sum_n x_d^{(n)} r_{nk}}{P_{dk}} - \frac{\sum_n r_{nk} - \sum_n x_d^{(n)} r_{nk}}{1 - P_{dk}} = 0$$

Solving for P_{dk}:

$$P_{dk} \sum_n r_{nk} - P_{dk} \sum_n x_d^{(n)} r_{nk} = \sum_n x_d^{(n)} r_{nk} - P_{dk} \sum_n x_d^{(n)} r_{nk}$$

Thus,

$$\hat{P}_{dk} = \frac{\sum_n x_d^{(n)} r_{nk}}{\sum_n r_{nk}}$$

We have the maximizing parameters for the expected log-joint

$$\arg\max_{\pi,P} \left\langle \sum_n \log P(x^{(n)}, s^{(n)} | \pi, P) \right\rangle_{q(\{s^{(n)}\})}$$

thus obtaining an iterative update for the parameters π and P in the M-step of EM.

(d) The Python code for the EM algorithm:

```python
from dataclasses import dataclass
from typing import List, Tuple

import matplotlib.pyplot as plt
import numpy as np
from scipy.special import logsumexp

from src.constants import DEFAULT_SEED


@dataclass
class Theta:
    """
    log_pi: the logarithm of the mixing proportions (1, k)
    log_p_matrix: the logarithm of the probability where the (i,j)th element is the probability that
                  pixel j takes value 1 under mixture component i (d, k)
    """

    log_pi: np.ndarray
    log_p_matrix: np.ndarray

    @property
    def pi(self):
        return np.exp(self.log_pi)

    @property
    def p_matrix(self):
        d, k = self.log_p_matrix.shape
        image_dimension = int(np.sqrt(d))
        return np.exp(self.log_p_matrix).reshape(image_dimension, image_dimension, -1)

    @property
    def log_one_minus_p_matrix(self) -> np.ndarray:
        """
        Compute log(1-P) where P=exp(log_p_matrix)
        :return: an array of the same shape as log_p_matrix (d, k)
        """
        log_of_one = np.zeros(self.log_p_matrix.shape)
        stacked_sum = np.stack((log_of_one, self.log_p_matrix))
        weights = np.ones(stacked_sum.shape)
        weights[1] = -1  # scale p matrix by -1 for subtraction
        return np.array(logsumexp(stacked_sum, b=weights, axis=0))

    def log_pi_repeated(self, n: int):
        """
        Repeats the log_pi vector n times along axis 0
        :param n: number of repetitions
        :return: an array of shape (n, k)
        """
        return np.repeat(self.log_pi, n, axis=0)


def _init_params(k: int, d: int, seed: int = DEFAULT_SEED) -> Theta:
    """
    Random initialisation of theta parameters (log_pi and log_p_matrix)
    :param k: Number of components
    :param d: Image dimension (number of pixels in a single image)
    :param seed: seed initialisation for random methods
    :return: theta: the parameters of the model
    """
    np.random.seed(seed)
    return Theta(
        log_pi=np.log(np.random.dirichlet(np.ones(k), size=1)),
        log_p_matrix=np.log(np.random.uniform(low=0, high=1, size=(d, k))),
    )


def _compute_log_component_p_x_i_given_theta(x: np.ndarray, theta: Theta) -> np.ndarray:
    """
    Compute the unweighted probability of each mixing component for each image
    :param x: the image data (n, d)
    :param theta: the parameters of the model
    :return: an array of the unweighted probabilities (n, k)
    """
    return x @ theta.log_p_matrix + (1 - x) @ theta.log_one_minus_p_matrix


def _compute_log_p_x_i_given_theta(x: np.ndarray, theta: Theta) -> np.ndarray:
    """
    Computes the log likelihood of each image in the dataset x
    :param x: the image data (n, d)
    :param theta: the parameters of the model
    :return: log_p_x_i_given_theta: a log likelihood array containing the log likelihood of each image (n
        ,1)
    """
    n, _ = x.shape
    log_component_probabilities = _compute_log_component_p_x_i_given_theta(
        x, theta
    )  # (n, k)
    return np.array(
        logsumexp(
            log_component_probabilities
            + theta.log_pi_repeated(n),  # scale each component by component probability
            axis=1,
```

```
 94              )
 95          )
 96
 97
 98      def _compute_log_likelihood(x: np.ndarray, theta: Theta) -> float:
 99          """
100          Computes the log likelihood of all images in the dataset x
101          :param x: the image data (n, d)
102          :param theta: the parameters of the model
103          :return: log_p_x_given_theta: the log likelihood array across all images
104          """
105          return np.sum(_compute_log_p_x_i_given_theta(x, theta)).item()
106
107
108      def _compute_log_e_step(x: np.ndarray, theta: Theta) -> np.ndarray:
109          """
110          Compute the e step of expectation maximisation
111          :param x: the image data (n, d)
112          :param theta: the parameters of the model
113          :return: an array of the log responsibilities of k mixture components for each image (n, k)
114          """
115          log_r_unnormalised = _compute_log_component_p_x_i_given_theta(x, theta)
116          log_r_normaliser = logsumexp(log_r_unnormalised, axis=1)
117          log_responsibility = log_r_unnormalised - log_r_normaliser[:, np.newaxis]
118          return log_responsibility
119
120
121      def _compute_log_pi_hat(log_responsibility: np.ndarray) -> np.ndarray:
122          """
123          Compute the log of the maximised mixing proportions
124          :param log_responsibility: an array of the log responsibilities of k mixture components for each image
               (n, k)
125          :return: an array of the maximised log mixing proportions (1, k)
126          """
127          n, _ = log_responsibility.shape
128          summed_responsibilities = logsumexp(log_responsibility, axis=0)
129
130          alpha = 2
131          beta = 2
132          summed_responsibilities = logsumexp(
133              np.stack(
134                  (
135                      (alpha - 1) * np.ones(summed_responsibilities.shape),
136                      summed_responsibilities,
137                  ),
138                  axis=0,
139              ),
140              axis=0,
141          )
142          return (summed_responsibilities - np.log(n + alpha + beta - 2)).reshape(1, -1)
143
144
145      def _compute_log_p_matrix_hat(
146          x: np.ndarray, log_responsibility: np.ndarray
147      ) -> np.ndarray:
148          """
149          Compute the log of the maximised pixel probabilities
150          :param x: the image data (n, d)
151          :param log_responsibility: an array of the log responsibilities of k mixture components for each image
               (n, k)
152          :return: an array of the maximised pixel probabilities for each component (d, k)
153          """
154          n, d = x.shape
155          _, k = log_responsibility.shape
156
157          x_repeated = np.repeat(x[:, :, np.newaxis], k, axis=2)  # (n, d, k)
158          log_responsibility_repeated = np.repeat(
159              log_responsibility[:, np.newaxis, :], d, axis=1
160          )  # (n, d, k)
161
162          log_p_matrix_unnormalised = logsumexp(
163              log_responsibility_repeated, b=x_repeated, axis=0
164          )  # (d, k)
165          log_p_matrix_normaliser = np.array(
166              logsumexp(log_responsibility_repeated, axis=0)
167          )  # (d, k)
168
169          alpha = 2
170          beta = 2
171          log_p_matrix_unnormalised = logsumexp(
172              np.stack(
173                  (
174                      (alpha - 1) * np.ones(log_p_matrix_unnormalised.shape),
175                      log_p_matrix_unnormalised,
176                  ),
177                  axis=0,
178              ),
179              axis=0,
180          )
181          log_p_matrix_normaliser = logsumexp(
182              np.stack(
183                  (
184                      (alpha + beta - 2) * np.ones(log_p_matrix_normaliser.shape),
185                      log_p_matrix_normaliser,
186                  ),
187                  axis=0,
```

```python
188            ),
189            axis=0,
190        )
191        log_p_matrix_normalised = log_p_matrix_unnormalised - log_p_matrix_normaliser
192        return log_p_matrix_normalised
193
194
195    def _compute_log_m_step(x: np.ndarray, log_responsibility: np.ndarray) -> Theta:
196        """
197        Compute the m step of expectation maximisation
198        :param x: the image data (n, d)
199        :param log_responsibility: an array of the log responsibilities of k mixture components for each image
            (n, k)
200        :return: thetas optimised after maximisation step
201        """
202        return Theta(
203            log_pi=_compute_log_pi_hat(log_responsibility),
204            log_p_matrix=_compute_log_p_matrix_hat(x, log_responsibility),
205        )
206
207
208    def _run_expectation_maximisation(
209        x: np.ndarray, theta: Theta, max_number_of_steps: int, epsilon: float
210    ) -> Tuple[Theta, List[float]]:
211        """
212        Run the expectation maximisation algorithm
213        :param x: the image data (n, d)
214        :param theta: initial theta parameters
215        :param max_number_of_steps: the maximum number of steps to run the algorithm
216        :param epsilon: the minimum required change in log likelihood, otherwise the algorithm stops early
217        :return: a tuple containing the optimised thetas and the log likelihood at each step of the algorithm
218        """
219        log_likelihoods = []
220        for _ in range(max_number_of_steps):
221            log_r = _compute_log_e_step(x, theta)
222            theta = _compute_log_m_step(x, log_r)
223
224            log_likelihoods.append(_compute_log_likelihood(x, theta))
225
226            # check for early stopping
227            if len(log_likelihoods) > 1:
228                if (log_likelihoods[-1] - log_likelihoods[-2]) < epsilon:
229                    break
230        return theta, log_likelihoods
231
232
233    def _plot_p_matrix(
234        thetas: List[Theta], ks: List[int], figure_title: str, figure_path: str
235    ):
236        n = len(ks)
237        m = np.max(ks)
238        fig = plt.figure()
239        for i, k in enumerate(ks):
240            for j in range(k):
241                ax = plt.subplot(n, m, m * i + j + 1)
242                ax.imshow(
243                    thetas[i].p_matrix[:, :, j],
244                    interpolation="None",
245                )
246                ax.tick_params(
247                    axis="x",
248                    which="both",
249                    bottom=False,
250                    top=False,
251                )
252                ax.tick_params(
253                    axis="y",
254                    which="both",
255                    left=False,
256                    right=False,
257                )
258                ax.xaxis.set_ticklabels([])
259                ax.yaxis.set_ticklabels([])
260                if j == 0:
261                    ax.set_ylabel(f"{k=}")
262        fig.suptitle(figure_title)
263        plt.savefig(figure_path)
264
265
266    def _plot_log_likelihoods(
267        log_likelihoods: List[List[float]],
268        ks: List[int],
269        figure_title: str,
270        figure_path: str,
271    ) -> None:
272        fig, ax = plt.subplots(len(ks), 1, constrained_layout=True)
273        fig.set_figwidth(10)
274        fig.set_figheight(15)
275        for i, k in enumerate(ks):
276            ax[i].plot(np.arange(1, len(log_likelihoods[i]) + 1), log_likelihoods[i])
277            ax[i].set_xlabel("Step")
278            ax[i].set_ylabel("Log-Likelihood")
279            ax[i].set_title(f"{k=}")
280        plt.suptitle(figure_title)
281
282        plt.savefig(figure_path)
```

```python
def e(
    x: np.ndarray,
    number_of_trials: int,
    ks: List[int],
    epsilon: float,
    max_number_of_steps: int,
    figure_path: str,
    figure_title: str,
) -> None:
    n, d = x.shape
    seeds = np.random.randint(
        low=number_of_trials * len(ks), size=(number_of_trials, len(ks))
    )
    for i in range(number_of_trials):
        init_thetas = []
        em_thetas = []
        log_likelihoods = []
        for j, k in enumerate(ks):
            init_theta = _init_params(k, d, seed=seeds[i, j])
            em_theta, log_likelihood = _run_expectation_maximisation(
                x,
                theta=init_theta,
                epsilon=epsilon,
                max_number_of_steps=max_number_of_steps,
            )
            init_thetas.append(init_theta)
            em_thetas.append(em_theta)
            log_likelihoods.append(log_likelihood)

        _plot_p_matrix(
            init_thetas,
            ks,
            figure_title=f"{figure_title} Trial {i}: Initialised P",
            figure_path=f"{figure_path}-{i}-initialised-p.png",
        )
        _plot_p_matrix(
            em_thetas,
            ks,
            figure_title=f"{figure_title} Trial {i}: EM Optimised P",
            figure_path=f"{figure_path}-{i}-optimised-p.png",
        )
        _plot_log_likelihoods(
            log_likelihoods,
            ks,
            figure_title=f"{figure_title} Trial {i}: Negative Log-Likelihood",
            figure_path=f"{figure_path}-{i}-log-like.png",
        )
```

src/solutions/q3.py

Run your algorithm on the data set for values of K in 2, 3, 4, 7, 10. Plot the log likelihood as a function of the iteration number, and display the parameters found. [30 marks]

(e) Run the algorithm a few times starting from randomly chosen initial conditions. Do you obtain the same solutions (up to permutation)? Does this depend on K? Show the learned probability vectors as images.

Comment on how well the algorithm works, whether it finds good clusters (look at the cluster means and responsibilities and try to interpret them), and how you might improve the model. [10 marks]

Question 5

(a) Let $p(\mathbf{s}_i, \mathbf{s}_{i-1})$ be the probability of the pair of symbols $\mathbf{s}_i$ and $\mathbf{s}_{i-1}$ occurring together in the text ($\mathbf{s}_{i-1}$ followed by $\mathbf{s}_i$ where order matters). We can model $p(\mathbf{s}_i, \mathbf{s}_{i-1})$ as a multinomial distribution with $N = 1$ and $D = 53^2$:

$$p(\mathbf{s}_i, \mathbf{s}_{i-1}) = \frac{1}{s^1! s^2! \cdots s^{53^2}} \prod_{j \in \{1,...53\}, k \in \{1,...53\}} p_{s_i, s_{i-1}}^{t^{s_i, s_{i-1}}}$$

where $t^{s_i, s_{i-1}}$ is an indicator of transition s_{i-1} to s_i. For convenience we will denote $t^{(\alpha, \beta)}$) as the transition from A multinomial distribution is appropriate because there can only be only one of 53 symbols chosen as s (i.e. a 53 dimensional dice). Thus, $p(s = s_i)$ represents the probability of the symbol s_i in the text.

We can convert p(s) into exponential family form:

$$p(\mathbf{s}) = \frac{1}{s^1! s^2! \cdots s^D} \exp\left(\mathbf{s}^T \log(\mathbf{p})\right)$$

where $\mathbf{p}$ is the vector of p_i's. Thus the sufficient statistic is $T(\mathbf{s}) = \mathbf{s}^T$.

The ML estimate

(b) The latent variables $\sigma(s)$ for different symbols s are not independent. This is because by choosing an encoding for one symbol $e = \sigma(s)$, the encoding for a second symbol $\sigma(s')$ cannot be e. We have 53 symbols but only 52 degrees of freedom, because once we have defined the encoding for 52 symbols, the encoding for the 53^{rd} symbol cannot be chosen. Thus, there exists a dependence between the symbols for a given σ.

The joint probability of the encrypted text $e_1 e_2 \cdots e_n$ given σ:

$$P(e_1, e_2, ..., e_n | \sigma) = \psi(\gamma = \sigma^{-1}(e_1)) \prod_{i=2}^{n} \psi(\alpha = \sigma^{-1}(e_i), \beta = \sigma^{-1}(e_{i-1}))$$

because σ is the encoding function, mapping to a symbol s into the encoded text as e, we require σ^{-1} the decoding function mapping the encoded symbol e back to s.

(c) The proposal probability $S(\sigma \to \sigma')$ depends on the permutations of σ and σ' because we only choose a proposal σ' that differs at *two* spots:

$$\sigma'(s^i) = \sigma(s^j)$$
$$\sigma'(s^j) = \sigma(s^i)$$

for any two symbols s^i and s^j of the 53 possible symbols ($s^i \neq s^j$).

Therefore, if the above doesn't hold for σ', $S(\sigma \to \sigma') = 0$, because with our method of choosing proposals, it is not possible to choose σ'. At σ there are $\binom{53}{2}$ possible proposal σ''s with the above property. Because we are assuming a uniform prior distribution over σ's, the transition probability of a σ' that satisfies the above property is $S(\sigma \to \sigma') = \frac{1}{\binom{53}{2}}$.

The MH acceptance probability is given as:

$$A(\sigma \to \sigma'|\mathcal{D}) = \min\{1, \frac{S(\sigma' \to \sigma)P(\sigma'|\mathcal{D})}{S(\sigma \to \sigma')P(\sigma|\mathcal{D})})\}$$

where $S(\sigma \to \sigma')$ is the conditional transition probability of σ' given σ and $\mathcal{D}$ is our encrypted text $e_1, e_2, ..., e_n$.

$S(\sigma \to \sigma') = S(\sigma' \to \sigma)$ for all σ and σ' that differ only at two spots because the probability in this case will always be $\frac{1}{\binom{53}{2}}$, we can simplify:

$$A(\sigma \to \sigma'|\mathcal{D}) = \min\{1, \frac{P(\sigma'|\mathcal{D})}{P(\sigma|\mathcal{D})})\}$$

From Bayes' Theorem:

$$P(\sigma|\mathcal{D}) = \frac{P(\mathcal{D}|\sigma)P(\sigma)}{\sum_{\sigma'} P(\mathcal{D}|\sigma')P(\sigma')}$$

We are assuming a uniform prior for σ, so $P(\sigma)$ is a constant and we can simplify further:

$$A(\sigma \to \sigma'|\mathcal{D}) = \min\{1, \frac{P(\mathcal{D}|\sigma')}{P(\mathcal{D}|\sigma)})\}$$

This is the acceptance probability for a given proposal σ'. The expression for $P(\mathcal{D}|\sigma)$ is $P(e_1, e_2, ..., e_n|\sigma)$ described in the previous part.

(d) The Python code for the MH sampler:

src/solutions/q5.py

Implement the MH sampler, and run it on the provided encrypted text. Report the current decryption of the first 60 symbols after every 100 iterations. Your Markov chain should converge to give you a fairly sensible message. (Hint: it may help to initialize your chain intelligently and to try multiple times; in any case, please describe what you did). [30 marks]

TODO

(e) Note that some $\Psi(\alpha, \beta)$ values may be zero. Does this affect the ergodicity of the chain? If the chain remains ergodic, give a proof; if not, explain and describe how you can restore ergodicity. [5 marks]

TODO

(f) Analyse this approach to decoding. For instance, would symbol probabilities alone (rather than transitions) be sufficient? If we used a second order Markov chain for English text, what problems might we encounter? Will it work if the encryption scheme allows two symbols to be mapped to the same encrypted value? Would it work for Chinese with > 10000 symbols? [13 marks]

TODO

Question 7

(a) To find the local extrema of the function $f(x, y) = x + 2y$ subject to the constraint $y^2 + xy = 1$, first we define $g(x, y)$:

$$g(x, y) = y^2 + xy - 1$$

where $g(x, y) = 0$ is an equivalent representation of the given constraint.

We can therefore construct the optimisation problem:

$$\min_{\mathbf{x}} f(\mathbf{x})$$

such that $g(\mathbf{x}) = \mathbf{0}$ and $\mathbf{x} := [x, y]^T$.

We can calculate $\nabla f(\mathbf{x})$:

$$\nabla f(\mathbf{x}) = [\frac{\partial}{\partial x}(x + 2y), \frac{\partial}{\partial y}(x + 2y)]^T$$

$$\nabla f(\mathbf{x}) = [1, 2]^T$$

and calculating $\nabla g(\mathbf{x})$:

$$\nabla g(\mathbf{x}) = [\frac{\partial}{\partial x}(y^2 + xy - 1), \frac{\partial}{\partial y}(y^2 + xy - 1)]^T$$

$$\nabla g(\mathbf{x}) = [y, 2y + x]^T$$

Solving the constraint optimisation problem with Lagrange multipliers, we set up the equations:

$$\nabla f(\mathbf{x}) + \lambda \nabla g(\mathbf{x}) = \mathbf{0}$$

and

$$g(\mathbf{x}) = 0$$

Giving us the three equations:

$$1 + \lambda y = 0$$
$$2 + \lambda(2y + x) = 0$$
$$y^2 + xy - 1 = 0$$

Substituting $\lambda y = -1$ from the first equation into the second equation:

$$2 + 2(-1) + x = 0$$

We see that $x = 0$. Solving for y in our third equation with $x = 0$:

$$y^2 - 1 = 0$$

We see that $y = \pm 1$ and from the first equation $\lambda \mp 1$.

The local extrema are $(x = 0, y = 1)$ when our $\lambda = -1$ and $(x = 0, y = -1)$ when our $\lambda = 1$.

(b)

    (i) Given that $g(a) = \ln(a)$, we want to transform this to the form $f(x, a) = 0$:

$$x = \ln(a)$$

$$\exp(x) - a = 0$$

    Thus,
$$f(x, a) = \exp(x) - a$$

    (ii) We know that for Newton's method's

$$x_{n+1} = x_n - \frac{f(x_n)}{f'(x_n)}$$

    where $f(x_n) = f(x_n, a) = \exp(x_n) - a$

    We can calculate:
$$f'(x) = \frac{\partial f(x, a)}{\partial x} = \exp(x)$$

    Assuming we can evaluate $\exp(x)$, our update equation:

$$x_{n+1} = x_n - \frac{\exp(x_n) - a}{\exp(x_n)}$$

    Simplifying:
$$x_{n+1} = x_n + \frac{a}{\exp(x_n)} - 1$$

Appendix: main.py

```python
import os

import numpy as np

from src.constants import BINARY_DIGITS_FILE_PATH, OUTPUTS_FOLDER
from src.solutions import q1, q2, q3

if __name__ == "__main__":

    if not os.path.exists(OUTPUTS_FOLDER):
        os.makedirs(OUTPUTS_FOLDER)

    x = np.loadtxt(BINARY_DIGITS_FILE_PATH)
    # Question 1
    Q1_OUTPUT_FOLDER = os.path.join(OUTPUTS_FOLDER, "q1")
    if not os.path.exists(Q1_OUTPUT_FOLDER):
        os.makedirs(Q1_OUTPUT_FOLDER)

    q1.d(
        x,
        figure_path=os.path.join(Q1_OUTPUT_FOLDER, "q1d.png"),
        figure_title="Q1d: Maximum Likelihood Estimate",
    )
    q1.e(
        x,
        alpha=3,
        beta=3,
        figure_path=os.path.join(Q1_OUTPUT_FOLDER, "q1e"),
        figure_title="Q1e: Maximum A Prior",
    )

    # Question 2
    Q2_OUTPUT_FOLDER = os.path.join(OUTPUTS_FOLDER, "q2")
    if not os.path.exists(Q2_OUTPUT_FOLDER):
        os.makedirs(Q2_OUTPUT_FOLDER)
    q2.c(x, table_path=os.path.join(Q2_OUTPUT_FOLDER, "q2c.csv"))

    # Question 3
    Q3_OUTPUT_FOLDER = os.path.join(OUTPUTS_FOLDER, "q3")
    if not os.path.exists(Q3_OUTPUT_FOLDER):
        os.makedirs(Q3_OUTPUT_FOLDER)
    q3.e(
        x,
        number_of_trials=3,
        ks=[2, 3, 4, 7, 10],
        epsilon=1e-1,
        max_number_of_steps=int(1e2),
        figure_path=os.path.join(Q3_OUTPUT_FOLDER, "q3e"),
        figure_title="Q3e",
    )
```

main.py