

COMP0086 Summative Assignment

Nov 14, 2022

Question 1

- (a) Our sample space for images is $\{0, 1\}^D$, where each of our D dimensions can only take binary values, D being the number of pixels in the image. The exponential family best suited on this sample space is the D -dimensional multivariate Bernoulli distribution because it shares the same sample space. On the other hand, a D -dimensional multivariate Gaussian has the sample space \mathbb{R}^D , which does not match the sample space of our data. To match our data sample space, we might have to define an additional mapping between our data and model spaces, adding unnecessary complexity. Thus it would be inappropriate to model this dataset of images with a multivariate Gaussian.
- (b) For $\{\mathbf{x}^{(n)}\}_{n=1}^N$, a data set of N images, the joint likelihood (assuming images are independently and identically distributed) is the product of N D -dimensional multivariate Bernoulli distributions, one for each image:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | \mathbf{p}) = \prod_{n=1}^N P(\mathbf{x}^{(n)} | \mathbf{p})$$

Substituting the D -dimensional multivariate Bernoulli:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | \mathbf{p}) = \prod_{n=1}^N \prod_{d=1}^D p_d^{x_d^{(n)}} (1 - p_d)^{1 - x_d^{(n)}}$$

Taking the logarithm of this, we get the log likelihood:

$$\mathcal{L}(\{\mathbf{x}^{(n)}\}_{n=1}^N | \mathbf{p}) = \sum_{n=1}^N \sum_{d=1}^D [x_d^{(n)} \log(p_d) + (1 - x_d^{(n)}) \log(1 - p_d)]$$

Note that since the logarithm is a monotonically increasing function on \mathbb{R}_+ , the maximisers and minimisers of the likelihood do not change. Thus, to solve for the maximum likelihood estimate, \hat{p}_d , we can take the derivative of $\mathcal{L}(\{\mathbf{x}^{(n)}\}_{n=1}^N | \mathbf{p})$ with respect to p_d , the d^{th} element of \mathbf{p} :

$$\begin{aligned} \frac{\partial \mathcal{L}(\{\mathbf{x}^{(n)}\}_{n=1}^N | \mathbf{p})}{\partial p_d} &= \sum_{n=1}^N \left(\frac{x_d^{(n)}}{p_d} - \frac{1 - x_d^{(n)}}{1 - p_d} \right) \\ \frac{\partial \mathcal{L}(\{\mathbf{x}^{(n)}\}_{n=1}^N | \mathbf{p})}{\partial p_d} &= \frac{\sum_{n=1}^N x_d^{(n)}}{p_d} - \frac{\sum_{n=1}^N (1 - x_d^{(n)})}{1 - p_d} \end{aligned}$$

and set the derivative to zero and solve for \hat{p}_d :

$$\begin{aligned}\frac{\sum_{n=1}^N x_d^{(n)}}{\hat{p}_d} - \frac{\sum_{n=1}^N (1 - x_d^{(n)})}{1 - \hat{p}_d} &= 0 \\ \sum_{n=1}^N x_d^{(n)} - \hat{p}_d \sum_{n=1}^N x_d^{(n)} - \hat{p}_d \cdot N + \hat{p}_d \sum_{n=1}^N x_d^{(n)} &= 0 \\ \hat{p}_d &= \frac{1}{N} \sum_{n=1}^N x_d^{(n)}\end{aligned}$$

Moreover, the second derivative with respect to p_d :

$$\frac{\partial \mathcal{L}(\{\mathbf{x}^{(n)}\}_{n=1}^N | \mathbf{p})}{\partial p_d^2} = \frac{-\sum_{n=1}^N x_d^{(n)}}{p_d^2} + \frac{-\sum_{n=1}^N (1 - x_d^{(n)})}{(1 - p_d)^2}$$

For a maximum, we need to show that the second derivative is negative. Since $x_d^{(n)} \in \{0, 1\}$, in the worst case, of $N = 1$, the single pixel $x_d^{(1)}$ must either be white ($\sum_{n=1}^N x_d^{(n)} > 0$) or black ($\sum_{n=1}^N 1 - x_d^{(n)} > 0$) with the other being zero, $\frac{\partial \mathcal{L}(\{\mathbf{x}^{(n)}\}_{n=1}^N | \mathbf{p})}{\partial p_d^2} < 0$ will be guaranteed and \hat{p}_d is a maximum as required for the maximum likelihood estimate.

Because we assume that each pixel is independent (we are taking the product of D one dimensional Bernoulli distributions), we can express the maximum likelihood for \mathbf{p} in vectorised form as $\hat{\mathbf{p}}^{MLE}$:

$$\hat{\mathbf{p}}^{MLE} = \frac{1}{N} \sum_{n=1}^N \mathbf{x}^{(n)}$$

(c) From Bayes' Theorem:

$$P(\mathbf{p} | \{\mathbf{x}^{(n)}\}_{n=1}^N) = \frac{P(\{\mathbf{x}^{(n)}\}_{n=1}^N | \mathbf{p}) P(\mathbf{p})}{P(\{\mathbf{x}^{(n)}\}_{n=1}^N)}$$

Taking the logarithm:

$$\mathcal{L}(\mathbf{p} | \{\mathbf{x}^{(n)}\}_{n=1}^N) = \mathcal{L}(\{\mathbf{x}^{(n)}\}_{n=1}^N | \mathbf{p}) + \mathcal{L}(\mathbf{p}) - \mathcal{L}(\{\mathbf{x}^{(n)}\}_{n=1}^N)$$

Taking the derivative with respect to p_d :

$$\frac{\partial \mathcal{L}(\mathbf{p} | \{\mathbf{x}^{(n)}\}_{n=1}^N)}{\partial p_d} = \frac{\partial \mathcal{L}(\{\mathbf{x}^{(n)}\}_{n=1}^N | \mathbf{p})}{\partial p_d} + \frac{\partial \mathcal{L}(\mathbf{p})}{\partial p_d}$$

where $\frac{\partial \mathcal{L}(\{\mathbf{x}^{(n)}\}_{n=1}^N)}{\partial p_d} = 0$ because it doesn't depend on p_d .

We know from (b):

$$\frac{\partial \mathcal{L}(\{\mathbf{x}^{(n)}\}_{n=1}^N | \mathbf{p})}{\partial p_d} = \frac{\sum_{n=1}^N x_d^{(n)}}{p_d} - \frac{\sum_{n=1}^N (1 - x_d^{(n)})}{1 - p_d}$$

For the second term $\frac{\partial \mathcal{L}(\mathbf{p})}{\partial p_d}$, we start with $P(\mathbf{p})$, assuming each pixel to have an independent prior:

$$P(\mathbf{p}) = \prod_{d=1}^D P(p_d)$$

Assuming a Beta prior on each p_d :

$$P(\mathbf{p}) = \prod_{d=1}^D \frac{1}{B(\alpha, \beta)} p_d^{\alpha-1} (1-p_d)^{\beta-1}$$

Taking the logarithm:

$$\mathcal{L}(\mathbf{p}) = \sum_{d=1}^D -\log(B(\alpha, \beta)) + (\alpha-1) \log p_d + (\beta-1) \log(1-p_d)$$

Taking the derivative with respect to p_d :

$$\frac{\partial \mathcal{L}(\mathbf{p})}{\partial p_d} = \frac{(\alpha-1)}{p_d} - \frac{(\beta-1)}{1-p_d}$$

Since we are only concerned with p_d , we are only left with a single element of the summation pertaining to p_d .

Combining, we have have an expression for $\frac{\partial \mathcal{L}(\mathbf{p}|\{\mathbf{x}^{(n)}\}_{n=1}^N)}{\partial p_d}$:

$$\frac{\partial \mathcal{L}(\mathbf{p}|\{\mathbf{x}^{(n)}\}_{n=1}^N)}{\partial p_d} = \frac{\sum_{n=1}^N x_d^{(n)}}{p_d} - \frac{\sum_{n=1}^N (1-x_d^{(n)})}{1-p_d} + \frac{(\alpha-1)}{p_d} - \frac{(\beta-1)}{1-p_d}$$

$$\frac{\partial \mathcal{L}(\mathbf{p}|\{\mathbf{x}^{(n)}\}_{n=1}^N)}{\partial p_d} = \frac{(\alpha-1) + \sum_{n=1}^N x_d^{(n)}}{p_d} - \frac{(\beta-1) + \sum_{n=1}^N (1-x_d^{(n)})}{1-p_d}$$

To find the maximum a posteriori (MAP) estimate \hat{p}_d set $\frac{\partial \mathcal{L}(\mathbf{p}|\{\mathbf{x}^{(n)}\}_{n=1}^N)}{\partial p_d} = 0$ and solve:

$$0 = \frac{(\alpha-1) + \sum_{n=1}^N x_d^{(n)}}{\hat{p}_d} - \frac{(\beta-1) + \sum_{n=1}^N (1-x_d^{(n)})}{1-\hat{p}_d}$$

$$0 = (1-\hat{p}_d)(\alpha-1) + (1-\hat{p}_d) \left(\sum_{n=1}^N x_d^{(n)} \right) - \hat{p}_d(\beta-1) - \hat{p}_d \left(\sum_{n=1}^N (1-x_d^{(n)}) \right)$$

$$0 = (\alpha - \alpha\hat{p}_d + \hat{p}_d - 1) + \left(\sum_{n=1}^N x_d^{(n)} - \hat{p}_d \sum_{n=1}^N x_d^{(n)} \right) - (\hat{p}_d\beta - \hat{p}_d) - \left(\hat{p}_d \cdot N - \hat{p}_d \sum_{n=1}^N x_d^{(n)} \right)$$

Cancelling the $\hat{p}_d \sum_{n=1}^N x_d^{(n)}$ terms:

$$0 = \alpha - \alpha\hat{p}_d + \hat{p}_d - 1 + \sum_{n=1}^N x_d^{(n)} - \hat{p}_d\beta + \hat{p}_d - \hat{p}_d \cdot N$$

$$0 = \hat{p}_d(2 - \alpha - \beta - N) + \alpha - 1 + \sum_{n=1}^N x_d^{(n)}$$

$$\hat{p}_d = \frac{\alpha - 1 + \sum_{n=1}^N x_d^{(n)}}{(N + \alpha + \beta - 2)}$$

To show that this is a maximum, the second derivative is:

$$\frac{\partial^2 \mathcal{L}(\mathbf{p} | \{\mathbf{x}^{(n)}\}_{n=1}^N)}{(\partial p_d)^2} = \frac{(1 - \alpha) - \sum_{n=1}^N x_d^{(n)}}{(p_d)^2} + \frac{(1 - \beta) - \sum_{n=1}^N (1 - x_d^{(n)})}{(1 - p_d)^2}$$

. For a maximum, we need $\frac{\partial^2 \mathcal{L}(\mathbf{p} | \{\mathbf{x}^{(n)}\}_{n=1}^N)}{(\partial p_d)^2} < 0$ meaning that we need at least one of the strict inequalities $\alpha < 1 - \sum_{n=1}^N x_d^{(n)}$ or $\beta < 1 - \sum_{n=1}^N (1 - x_d^{(n)})$ to be satisfied, where the other can be \leq . The Beta distribution requires $\alpha > 0$ and $\beta > 0$ so this requirement will always be satisfied (in the worst case of a single image, either $x_d^{(1)} = 1$ or $1 - x_d^{(1)} = 1$).

Due to independence of our likelihood and priors for each dimension, we can express the maximum a priori for \mathbf{p} in vectorised form as $\hat{\mathbf{p}}^{MAP}$:

$$\hat{\mathbf{p}}^{MAP} = \frac{\alpha - 1 + \sum_{n=1}^N \mathbf{x}^n}{(N + \alpha + \beta - 2)}$$

(d&e) The Python code for MLE and MAP:

```
1 import matplotlib.pyplot as plt
2 import numpy as np
3
4
5 def _compute_maximum_likelihood_estimate(x: np.ndarray) -> np.ndarray:
6     """
7     X: numpy array of shape (N, D)
8     """
9     return np.mean(x, axis=0)
10
11
12 def _compute_maximum_a_priori_estimate(
13     x: np.ndarray, alpha: float, beta: float
14 ) -> np.ndarray:
15     """
16     X: numpy array of shape (N, D)
17     alpha: param of prior distribution
18     beta: param of prior distribution
19     """
20
21     n, _ = x.shape
22     return (alpha - 1 + np.sum(x, axis=0)) / (n + alpha + beta - 2)
23
24
25 def d(x, figure_path, figure_title):
26     maximum_likelihood = _compute_maximum_likelihood_estimate(x)
27     plt.figure()
28     plt.imshow(
29         np.reshape(maximum_likelihood, (8, 8)),
30         interpolation="None",
31     )
32     plt.colorbar()
33     plt.axis("off")
34     plt.title(figure_title)
35     plt.savefig(figure_path)
36
37
38 def e(x, alpha, beta, figure_path, figure_title):
39     maximum_a_priori = _compute_maximum_a_priori_estimate(x, alpha, beta)
40     plt.figure()
41     plt.imshow(
42         np.reshape(maximum_a_priori, (8, 8)),
43         interpolation="None",
44     )
45     plt.colorbar()
46     plt.axis("off")
47     plt.title(figure_title)
48     plt.savefig(f"{figure_path}.png")
49
50     maximum_likelihood = _compute_maximum_likelihood_estimate(x)
51     plt.figure()
52     plt.imshow(
53         np.reshape(maximum_a_priori - maximum_likelihood, (8, 8)),
54         interpolation="None",
55     )
56     plt.colorbar()
57     plt.axis("off")
58     plt.title(f"MAP vs MLE")
59     plt.savefig(f"{figure_path}-mle-vs-map.png")
```

src/solutions/q1.py

Displaying the learned parameters:

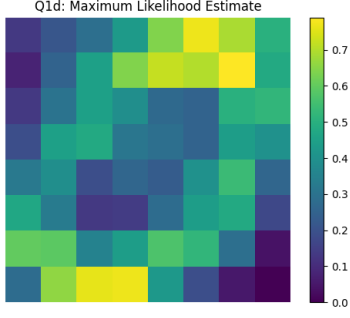


Figure 1: ML parameters

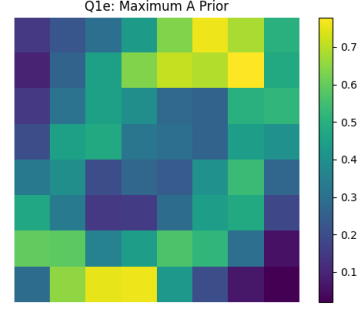


Figure 2: MAP parameters

Comparing the equations:

$$\hat{\mathbf{p}}^{MLE} = \frac{1}{N} \sum_{n=1}^N \mathbf{x}^{(n)}$$

and

$$\hat{\mathbf{p}}^{MAP} = \frac{\alpha - 1 + \sum_{n=1}^N \mathbf{x}^n}{(N + \alpha + \beta - 2)}$$

As the number of data points increases, $\hat{\mathbf{p}}^{MAP}$ approaches $\frac{1}{N} \sum_{n=1}^N \mathbf{x}^{(n)}$, $\hat{\mathbf{p}}^{MLE}$. This makes sense because as our data set gets bigger, we are less reliant on our prior. However, if a specific pixel in all of the images of our data set are white or all black, the MLE for that pixel would either be 1 or 0. This may not be representative of our intuitions about images, as there should be some non-zero probability of a pixel being black or white. By introducing an appropriate prior we can ensure that the probability of that pixel will never be exactly zero or one. In our case, with a Beta(3,3) prior on each pixel, our parameter values are biased to be closer to 0.5 and to never be at the extremities 0 and 1. We can see this in Figure 2 where the range of our parameters is smaller than the range of Figure 1 and doesn't include zero. Figure 3 visualises $\hat{\mathbf{p}}^{MAP} - \hat{\mathbf{p}}^{MLE}$ and we can see that for likelihoods greater than 0.5 in the MLE, the MAP has a lower value and for likelihoods less than 0.5, the MAP has a higher value, confirming our intuitions.

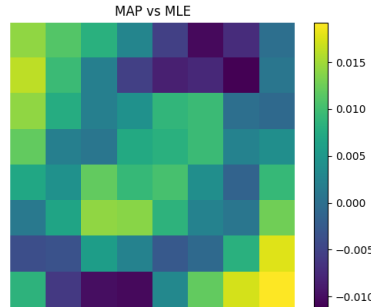


Figure 3: $\hat{\mathbf{p}}^{MAP} - \hat{\mathbf{p}}^{MLE}$

Priors can also help ensure numerical stability during calculations. The logarithm of zero is negative infinity, so having if the MLE is zero it can be problematic for log-likelihoods calculations whereas MAP can ensure non-zero probabilities. Interestingly, when $\alpha = \beta = 1$, $\hat{\mathbf{p}}^{MLE} = \hat{\mathbf{p}}^{MAP}$. This is when the prior is a uniform distribution and so there is uniform bias on the location of \mathbf{p} and we recover the MLE.

On the other hand, a mis-specified prior can be problematic, as the estimated parameters might be skewed by the prior and not properly represent the underlying data generating process, this can result in parameter estimates that are worse than using the MLE if our data set is limited.

Question 2

When all D components are generated from a Bernoulli distribution with $p_d = 0.5$, we have the likelihood function for model M_1 :

$$P(\mathbf{x}^{(n)}|\mathbf{p}^{(1)} = [0.5, 0.5, \dots, 0.5]^T, M_1) = \prod_{n=1}^N \prod_{d=1}^D (0.5)^{x_d^{(n)}} (0.5)^{1-x_d^{(n)}}$$

When all D components are generated from Bernoulli distributions with unknown, but identical, p_d , we have the likelihood function for model M_2 :

$$P(\mathbf{x}^{(n)}|\mathbf{p}^{(2)} = [p_d, p_d, \dots, p_d]^T, M_2) = \prod_{n=1}^N \prod_{d'=1}^D p_d^{x_{d'}^{(n)}} (1 - p_d)^{1-x_{d'}^{(n)}}$$

When each component is Bernoulli distributed with separate, unknown p_d , we have the likelihood function for model M_3 :

$$P(\mathbf{x}^{(n)}|\mathbf{p}^{(3)} = [p_1, p_2, \dots, p_D]^T, M_3) = \prod_{n=1}^N \prod_{d=1}^D p_d^{x_d^{(n)}} (1 - p_d)^{1-x_d^{(n)}}$$

For each model M_i , we can marginalise out $\mathbf{p}^{(i)}$ to get $P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_i)$:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_i) = \int_0^1 \dots \int_0^1 P(\{\mathbf{x}^{(n)}\}_{n=1}^N | p_d, M_i) P(p_d | M_i) dp_1 \dots dp_D$$

where $d = 1, \dots, D$ and $\{\mathbf{x}^{(n)}\}_{n=1}^N$ is our data set.

Given that the prior of any unknown probabilities is uniform, i.e. $P(p_d | M_i) = 1$. We can simplify:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_i) = \int_0^1 \dots \int_0^1 P(\{\mathbf{x}^{(n)}\}_{n=1}^N | p_d, M_i) dp_1 \dots dp_D$$

For M_1 , we have that all pixels have probability 0.5:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_1) = \int_0^1 \dots \int_0^1 \prod_{n=1}^N \prod_{d=1}^D (0.5)^{x_d^{(n)}} (1 - 0.5)^{1-x_d^{(n)}} d\theta_1 \dots d\theta_D$$

We can remove the integrals and knowing that either $x_d^{(n)}$ or $1 - x_d^{(n)}$ will be 1 and the other zero, we can simplify $(0.5)^{x_d^{(n)}} (1 - 0.5)^{1-x_d^{(n)}}$ to 0.5:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_1) = \prod_{n=1}^N \prod_{d=1}^D (0.5)$$

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_1) = (0.5)^{N \cdot D}$$

For M_2 , we have that all pixels share some probability p_d so we only need to integrate over a single variable p_d :

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_2) = \int_0^1 \prod_{n=1}^N \prod_{d'=1}^D p_d^{x_{d'}^{(n)}} (1 - p_d)^{1-x_{d'}^{(n)}} dp_d$$

Changing the products to sums:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_2) = \int_0^1 p_d^{\sum_{n=1}^N \sum_{d'=1}^D x_{d'}^{(n)}} (1 - p_d)^{\sum_{j=1}^N \sum_{d'=1}^D 1 - x_{d'}^{(n)}} dp_d$$

Rewriting:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_2) = \int_0^1 (p_d)^K (1 - p_{d'=1})^{N \cdot D - K} dp_d$$

where $K = \sum_{n=1}^N \sum_{d'=1}^D x_{d'}^{(n)}$.

This integral is the beta function:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_2) = \frac{K!(N \cdot D - K)!}{(N \cdot D + 1)!}$$

For M_3 , we need an integral for each p_d :

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_3) = \int_0^1 \dots \int_0^1 \prod_{n=1}^N \prod_{d=1}^D p_d^{x_d^{(n)}} (1 - p_d)^{1 - x_d^{(n)}} dp_1 \dots dp_D$$

We can separate the integrals to only contain the relevant p_d :

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_3) = \prod_{d=1}^D \left(\int_0^1 \prod_{n=1}^N p_d^{x_d^{(n)}} (1 - p_d)^{1 - x_d^{(n)}} dp_d \right)$$

Changing the products to sums:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_3) = \prod_{d=1}^D \left(\int_0^1 p_d^{\sum_{n=1}^N x_d^{(n)}} (1 - p_d)^{\sum_{n=1}^N 1 - x_d^{(n)}} dp_d \right)$$

In this case, we have the product of integrals where each evaluates to a beta function:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_3) = \prod_{d=1}^D \frac{K_d!(N - K_d)!}{(N + 1)!}$$

where $K_d = \sum_{n=1}^N x_d^{(n)}$.

The posterior probability of a model M_i can be expressed:

$$P(M_i | \{\mathbf{x}^{(n)}\}_{n=1}^N) = \frac{P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_i) P(M_i)}{P(\{\mathbf{x}^{(n)}\}_{n=1}^N)}$$

We only have three models, so in this case the normalisation $P(\{\mathbf{x}^{(n)}\}_{n=1}^N)$ can be expressed as a sum:

$$P(M_i | \{\mathbf{x}^{(n)}\}_{n=1}^N) = \frac{P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_i) P(M_i)}{\sum_{i \in \{1, 2, 3\}} P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_i) P(M_i)}$$

Given that $P(M_i) = \frac{1}{3}$ for all $i \in \{1, 2, 3\}$:

$$P(M_i | \{\mathbf{x}^{(n)}\}_{n=1}^N) = \frac{P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_i)}{\sum_{i \in \{1, 2, 3\}} P(\{\mathbf{x}^{(n)}\}_{n=1}^N | M_i)}$$

i	$P(M_i \{\mathbf{x}^{(n)}\}_{n=1}^N)$
1	1E-1924
2	1E-1858
3	$1-(1\text{E-}1924)-(1\text{E-}1858)$

Table 1: Posterior Probabilities

Calculating the posterior probabilities of each of the three models having generated the data in binarydigits.txt using python, we can show the values in the Table 1:

We can see that for models specified to have the same parameter value for all pixels like M_1 is very unlikely with the given data set. This makes sense because it is specifying models where the image is essentially blank (a uniform shade), which is not reflective of our digit images. Moreover, M_1 specifies a specific value of 0.5 for all the parameters whereas M_2 specifies any value for all the parameters as long as it's the same. So the model M_1 is a subset of the models specified in M_2 and we can see this reflected in our probabilities when $P(M_2|\{\mathbf{x}^{(n)}\}_{n=1}^N) > P(M_1|\{\mathbf{x}^{(n)}\}_{n=1}^N)$.

The Python code for calculating the posterior probabilities of the three models:

```
1 import numpy as np
2 import pandas as pd
3 from scipy.special import betaln, logsumexp
4
5
6 def _log-p-d-given-m1(x):
7     n, d = x.shape
8     return n * d * np.log(0.5)
9
10
11 def _log-p-d-given-m2(x):
12     n, d = x.shape
13     k = np.sum(x, axis=0).astype(int)
14     return betaln(np.sum(k) + 1, n * d - np.sum(k) + 1)
15
16
17 def _log-p-d-given-m3(x):
18     n, _ = x.shape
19     k = np.sum(x, axis=0).astype(int)
20     return logsumexp(betaln(k + 1, n - k + 1))
21
22
23 def c(x, table_path):
24     log-p-d-given-m = np.array(
25         [
26             _log-p-d-given-m1(x),
27             _log-p-d-given-m2(x),
28             _log-p-d-given-m3(x),
29         ]
30     )
31     log-p-m-given-d = log-p-d-given-m - logsumexp(log-p-d-given-m)
32     df = pd.DataFrame(
33         data=np.array(
34             [
35                 np.arange(len(log-p-m-given-d)).astype(int) + 1,
36                 [f"1E{int(x/np.log(10))}" for x in log-p-m-given-d[: -1]]
37                 + [
38                     f"1-{'-'.join([f'(1E{int(x/np.log(10))})' for x in log-p-m-given-d[: -1]])}"
39                 ],
40             ],
41         ).T,
42         columns=["Model", "P(M_i|D)"],
43     )
44     df.set_index("Model", inplace=True)
45     df.to_csv(table_path)
```

src/solutions/q2.py

Question 3

- (a) The likelihood for a model consisting of a mixture of K multivariate Bernoulli distributions can be expressed as the product across N data points:

$$P(\{\mathbf{x}^{(n)}\}_{n=1}^N|\theta) = \prod_{i=1}^N P(x_i|\theta)$$

where $\{\mathbf{x}^{(n)}\}_{n=1}^N$ is our data set with $\mathbf{x}^{(n)} \in \mathbb{R}^{D \times 1}$ and $\theta = \{\pi, \mathbf{P}\}$, $\pi = [\pi_1, \dots, \pi_K] \in \mathbb{R}^{K \times 1}$ our mixing proportions ($0 \leq \pi_k \leq 1$; $\sum_k \pi_k = 1$) and $\mathbf{P} \in \mathbb{R}^{D \times K}$ the K Bernoulli parameter vectors with elements p_{kd} denoting the probability that pixel d takes value 1 under mixture component k . We also assume the images are iid and that the pixels are independent of each other within each component distribution.

For each $P(\mathbf{x}^{(n)}|\theta)$:

$$P(\mathbf{x}^{(n)}|\theta) = \sum_{k=1}^K \pi_k \prod_{d=1}^D (p_{kd})^{\mathbf{x}_d^{(n)}} (1 - p_{kd})^{1 - \mathbf{x}_d^{(n)}}$$

The log-likelihood $\mathcal{L}(\mathbf{x}^{(n)}|\theta)$ can be expressed in matrix form:

$$\mathcal{L}(\mathbf{x}^{(n)}|\theta) = \log \sum_{k=1}^K \pi_k \exp \left(\mathbf{x}^{(n)} \log(\mathbf{P}_k) + (1 - \mathbf{x}^{(n)}) \log(1 - \mathbf{P}_k) \right)$$

which can be further vectorised using Python scipy's *logsumexp* operation.

Moreover, the log-likelihood $\mathcal{L}(\{\mathbf{x}^{(n)}\}_{n=1}^N|\theta)$ can be expressed:

$$\mathcal{L}(\{\mathbf{x}^{(n)}\}_{n=1}^N|\theta) = \sum_{i=1}^N \left(\log \sum_{k=1}^K \pi_k \exp \left(\mathbf{x}^{(n)} \log(\mathbf{P}_k) + (1 - \mathbf{x}^{(n)}) \log(1 - \mathbf{P}_k) \right) \right)$$

- (b) We know that:

$$P(A|B) \propto P(B|A)P(A)$$

Thus,

$$P(s^{(n)} = k|\mathbf{x}^{(n)}, \pi, \mathbf{P}) \propto P(\mathbf{x}^{(n)}|s^{(n)} = k, \pi, \mathbf{P})P(s^{(n)} = k|\pi, \mathbf{P})$$

where $s^{(n)} \in \{1, \dots, K\}$ a discrete hidden variable with $P(s^{(n)} = k|\mathbf{x}^{(n)}, \pi) = \pi_k$. Note that $P(s^{(n)} = k|\mathbf{x}^{(n)}, \pi) = P(s^{(n)} = k|\mathbf{x}^{(n)}, \pi, \mathbf{P})$ as $s^{(n)}$ isn't dependent on \mathbf{P} .

Let $P(s^{(n)} = k|\mathbf{x}^{(n)}, \pi, \mathbf{P}) \propto P(s^{(n)})$ be the unnormalised responsibility \tilde{r}_{nk} . Using the mixture for component k , π_k and the likelihood function of component k :

$$\tilde{r}_{nk} = \pi_k \prod_{d=1}^D (p_{kd})^{\mathbf{x}_d^{(n)}} (1 - p_{kd})^{1 - \mathbf{x}_d^{(n)}}$$

Normalising across the components:

$$r_{nk} = \frac{\tilde{r}_{nk}}{\sum_{j=1}^K \tilde{r}_{nj}}$$

we have calculated $P(s^{(n)} = k | \mathbf{x}^{(n)}, \pi, \mathbf{P})$ for the E step of an EM algorithm.

Moreover,

$$\log \tilde{r}_{nk} = \log \pi_k + \sum_{d=1}^D \left(\mathbf{x}_d^{(n)} \log(p_{kd}) + (1 - \mathbf{x}_d^{(n)}) \log(1 - \exp(\log(p_{kd}))) \right)$$

and

$$\log r_{nk} = \log \tilde{r}_{nk} - \log \sum_{j=1}^K \exp(\log \tilde{r}_{nj})$$

which can be vectorised as $\log \mathbf{r}_n$ calculated with $\log \pi$ and $\log \mathbf{P}$ using Python scipy's *logsumexp* operation.

(c) We know that the expectation log joint can be expressed:

$$\left\langle \sum_n \log P(\mathbf{x}^{(n)}, s^{(n)} | \pi, \mathbf{P}) \right\rangle_{q(\{s^{(n)}\})} = \sum_{n=1}^N q(s^{(n)}) \log P(\mathbf{x}^{(n)}, s^{(n)} | \pi, \mathbf{P})$$

Let this quantity be E . Each term of E can be expressed:

$$q(s^{(n)}) = \mathbf{r}_n$$

and

$$\log P(\mathbf{x}^{(n)}, s^{(n)} | \pi, \mathbf{P}) = \log[P(\mathbf{x}^{(n)} | s^{(n)}, \pi, \mathbf{P}) P(s^{(n)} | \pi, \mathbf{P})]$$

which is the vectorised version of $\log \tilde{r}_{nk}$ from part (b) so:

$$\log P(\mathbf{x}^{(n)}, s^{(n)} | \pi, \mathbf{P}) = \log(\pi) + \log(\mathbf{P})^T \mathbf{x}^{(n)} + \log(1 - \mathbf{P})^T (1 - \mathbf{x}^{(n)})$$

Combining:

$$E = \sum_n \mathbf{r}_n^T [\log(\pi) + \log(\mathbf{P})^T \mathbf{x}^{(n)} + \log(1 - \mathbf{P})^T (1 - \mathbf{x}^{(n)})]$$

To maximise with respect to π and \mathbf{P} for the M step, we want to take the derivative, set to zero, and solve for $\hat{\pi}$ and $\hat{\mathbf{P}}$.

For the k^{th} element of π :

$$\frac{\partial E}{\partial \pi_k} = \sum_n r_{nk} \frac{1}{\pi_k}$$

The second derivative:

$$\frac{\partial E}{(\partial \pi_k)^2} = \sum_n r_{nk} \frac{-1}{(\pi_k)^2}$$

is always negative because $r_{nk} \geq 0$, $\sum_n r_{nk} = 1$, $\pi_k \geq 0$, and $\sum_n \pi_k = 1$, ensuring a maximum in the next step.

We can calculate the maximiser with:

$$\frac{\partial E}{\partial \pi_k} + \lambda = 0$$

where λ is a Lagrange multiplier ensuring that the mixing proportions sum to unity.

Thus,

$$\hat{\pi}_k = \frac{\sum_n r_{nk}}{N}$$

For the dk^{th} element of \mathbf{P} :

$$\frac{\partial E}{\partial \mathbf{P}_{dk}} = \sum_n r_{nk} \frac{\partial}{\partial \mathbf{P}_{dk}} [\mathbf{x}_d^{(n)} \log \mathbf{P}_{dk} + (1 - \mathbf{x}_d^{(n)}) \log(1 - \mathbf{P}_{dk})]$$

Simplifying:

$$\frac{\partial E}{\partial \mathbf{P}_{dk}} = \sum_n r_{nk} \left(\frac{\mathbf{x}_d^{(n)}}{\mathbf{P}_{dk}} - \frac{1 - \mathbf{x}_d^{(n)}}{1 - \mathbf{P}_{dk}} \right)$$

Similar to Question 1, we can see that taking second derivative, the term in the brackets will always be less than zero and with $r_{nk} \geq 0$ and $\sum_n r_{nk} = 1$, the second derivative will always be negative. This ensures that we have a maximum in the next step.

Setting the derivative to zero:

$$\frac{\sum_n \mathbf{x}_d^{(n)} r_{nk}}{\mathbf{P}_{dk}} - \frac{\sum_n r_{nk} - \sum_n \mathbf{x}_d^{(n)} r_{nk}}{1 - \mathbf{P}_{dk}} = 0$$

Solving for $\hat{\mathbf{P}}_{dk}$:

$$\hat{\mathbf{P}}_{dk} \sum_n r_{nk} - \hat{\mathbf{P}}_{dk} \sum_n \mathbf{x}_d^{(n)} r_{nk} = \sum_n \mathbf{x}_d^{(n)} r_{nk} - \hat{\mathbf{P}}_{dk} \sum_n \mathbf{x}_d^{(n)} r_{nk}$$

Thus,

$$\hat{\mathbf{P}}_{dk} = \frac{\sum_n \mathbf{x}_d^{(n)} r_{nk}}{\sum_n r_{nk}}$$

We have the maximizing parameters for the expected log-joint

$$\arg \max_{\pi, \mathbf{P}} \left\langle \sum_n \log P(\mathbf{x}^{(n)}, s^{(n)} | \pi, \mathbf{P}) \right\rangle_{q(\{s^{(n)}\})}$$

thus obtaining an iterative update for the parameters π and \mathbf{P} in the M-step of EM. For numerical stability, we can compute the maximisation step for the MAP of $\mathbf{P}, \hat{\mathbf{P}}_{dk}^{MAP}$ by solving:

$$\frac{\partial E'}{\partial \mathbf{P}_{dk}} = 0$$

where

$$E' = \sum_{n=1}^N q(s^{(n)}) \log P(\mathbf{P} | \pi, \mathbf{x}^{(n)}, s^{(n)})$$

and from Bayes':

$$\log P(\mathbf{P} | \pi, \mathbf{x}^{(n)}, s^{(n)}) = \log P(\mathbf{x}^{(n)}, s^{(n)} | \pi, \mathbf{P}) + \log P(\mathbf{P}) - \log P(\mathbf{x}^{(n)}, s^{(n)} | \pi)$$

Assuming an independent Beta prior on each pixel of each component:

$$\log P(\mathbf{P}) = \sum_{k=1}^K \sum_{d=1}^D -\log(B(\alpha, \beta)) + (\alpha - 1) \log \mathbf{P}_{dk} + (\beta - 1) \log(1 - \mathbf{P}_{dk})$$

and

$$\frac{\partial \log P(\mathbf{P})}{\partial \mathbf{P}_{dk}} = \frac{(\alpha - 1)}{\mathbf{P}_{dk}} - \frac{(\beta - 1)}{1 - \mathbf{P}_{dk}}$$

Thus, the derivative can be expressed as:

$$\frac{\partial E'}{\partial \mathbf{P}_{dk}} = \sum_n \left(r_{nk} \left(\frac{\partial \log P(\mathbf{x}^{(n)}, s^{(n)} | \pi, \mathbf{P})}{\partial \mathbf{P}_{dk}} + \frac{\partial \log P(\mathbf{P})}{\partial \mathbf{P}_{dk}} \right) \right)$$

Substituting the appropriate expressions:

$$\frac{\partial E'}{\partial \mathbf{P}_{dk}} = \sum_n \left(r_{nk} \left(\frac{\mathbf{x}_d^{(n)}}{\mathbf{P}_{dk}} - \frac{1 - \mathbf{x}_d^{(n)}}{1 - \mathbf{P}_{dk}} + \frac{(\alpha - 1)}{\mathbf{P}_{dk}} - \frac{(\beta - 1)}{1 - \mathbf{P}_{dk}} \right) \right)$$

Simplifying:

$$\frac{\partial E'}{\partial \mathbf{P}_{dk}} = \frac{\sum_n r_{nk}(\alpha - 1 + \mathbf{x}_d^{(n)})}{\mathbf{P}_{dk}} - \frac{\sum_n r_{nk}(\beta - \mathbf{x}_d^{(n)})}{1 - \mathbf{P}_{dk}}$$

For a maximum, we see that we need $\alpha > \mathbf{x}_d^{(n)} - 1$ or $\beta < \mathbf{x}_d^{(n)}$, both of which are satisfied knowing that $\alpha > 0$ and $\beta > 0$. Setting $\frac{\partial E'}{\partial \mathbf{P}_{dk}} = 0$ we can calculate $\hat{\mathbf{P}}_{dk}^{MAP}$:

$$\sum_n r_{nk}(\alpha - 1 + \mathbf{x}_d^{(n)}) - \hat{\mathbf{P}}_{dk} \sum_n r_{nk}(\alpha - 1 + \mathbf{x}_d^{(n)}) = \hat{\mathbf{P}}_{dk} \sum_n r_{nk}(\beta - \mathbf{x}_d^{(n)})$$

$$\hat{\mathbf{P}}_{dk}^{MAP} = \frac{\sum_n r_{nk}(\mathbf{x}_d^{(n)} + \alpha - 1)}{(\alpha + \beta - 1)(\sum_n r_{nk})}$$

As a sense check, we can see when setting $\alpha = 1$ and $\beta = 1$ we recover $\hat{\mathbf{P}}_{dk}^{MLE}$ as we would expect.

(d) Plotting the posterior likelihood as a function of the iteration number:

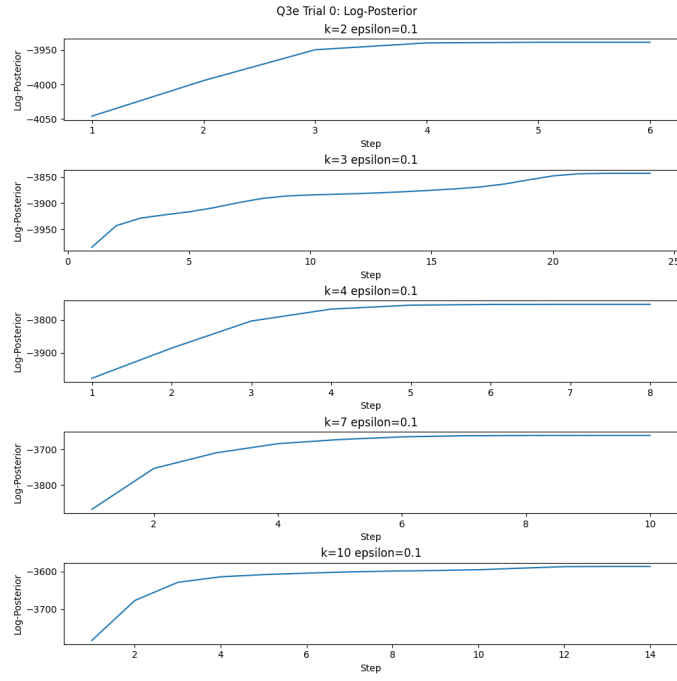


Figure 4: Log Likelihood vs Iteration Number

where *epsilon* is the stopping condition for the posterior posterior converges.

Displaying the parameters found for K in $\{2, 3, 4, 7, 10\}$:

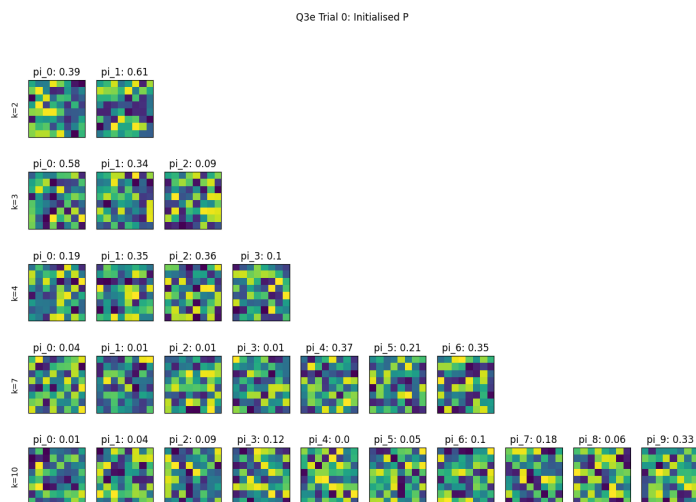


Figure 5: Randomly initialised parameters

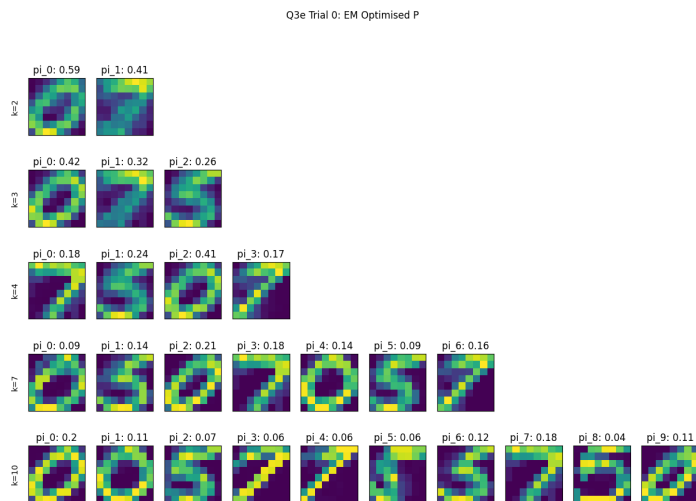


Figure 6: EM optimised parameters

The Python code for the EM algorithm:

```

1 from dataclasses import dataclass
2 from typing import List, Tuple
3
4 import matplotlib.pyplot as plt
5 import numpy as np
6 from scipy.special import logsumexp
7 from sklearn.manifold import TSNE
8
9 from src.constants import DEFAULT_SEED
10
11
12 @dataclass
13 class Theta:
14     """
15     log-pi: the logarithm of the mixing proportions (1, k)
16     log-p-matrix: the logarithm of the probability where the (i,j)th element is the probability that
17                   pixel j takes value 1 under mixture component i (d, k)
18     """
19
20     log-pi: np.ndarray
21     log-p-matrix: np.ndarray
22
23     @property
24     def pi(self):
25         return np.exp(self.log-pi)
26
27     @property
28     def p-matrix(self):
29         d, k = self.log-p-matrix.shape
30         image-dimension = int(np.sqrt(d))
31         return np.exp(self.log-p-matrix).reshape(image-dimension, image-dimension, -1)
32
33     @property
34     def log-one-minus-p-matrix(self) -> np.ndarray:
35         """
36         Compute log(1-P) where P=exp(log-p-matrix)
37         :return: an array of the same shape as log-p-matrix (d, k)
38         """
39         log-of-one = np.zeros(self.log-p-matrix.shape)
40         stacked_sum = np.stack((log-of-one, self.log-p-matrix))
41         weights = np.ones(stacked_sum.shape)
42         weights[1] = -1 # scale p matrix by -1 for subtraction
43         return np.array(logsumexp(stacked_sum, b=weights, axis=0))
44
45     def log-pi-repeated(self, n: int):
46         """
47         Repeats the log-pi vector n times along axis 0
48         :param n: number of repetitions
49         :return: an array of shape (n, k)
50         """
51         return np.repeat(self.log-pi, n, axis=0)
52
53
54 def _init_params(k: int, d: int) -> Theta:
55     """
56     Random initialisation of theta parameters (log-pi and log-p-matrix)
57     :param k: Number of components
58     :param d: Image dimension (number of pixels in a single image)
59     :return: theta: the parameters of the model
60     """
61     return Theta(
62         log-pi=np.log(np.random.dirichlet(np.ones(k), size=1)),
63         log-p-matrix=np.log(np.random.uniform(low=0, high=1, size=(d, k))),
64     )
65
66
67 def _compute_log-component-p-x-i-given-theta(x: np.ndarray, theta: Theta) -> np.ndarray:
68     """
69     Compute the unweighted probability of each mixing component for each image
70     :param x: the image data (n, d)
71     :param theta: the parameters of the model
72     :return: an array of the unweighted probabilities (n, k)
73     """
74     return x @ theta.log-p-matrix + (1 - x) @ theta.log-one-minus-p-matrix
75
76
77 def _compute_log-p-x-i-given-theta(x: np.ndarray, theta: Theta) -> np.ndarray:
78     """
79     Computes the log likelihood of each image in the dataset x
80     :param x: the image data (n, d)
81     :param theta: the parameters of the model
82     :return: log-p-x-i-given-theta: a log likelihood array containing the log likelihood of each image (n
83             ,1)
84     """
85     n, _ = x.shape
86     log-component-probabilities = _compute_log-component-p-x-i-given-theta(
87         x, theta
88     ) # (n, k)
89     return np.array(
90         logsumexp(
91             log-component-probabilities
92             + theta.log-pi-repeated(n), # scale each component by component probability
93             axis=1,
94         )
95 
```

```

94 )
95
96
97 def _compute_log_likelihood(x: np.ndarray, theta: Theta) -> float:
98     """
99     Computes the log likelihood of all images in the dataset x
100     :param x: the image data (n, d)
101     :param theta: the parameters of the model
102     :return: log_p_x_given_theta: the log likelihood array across all images
103     """
104     return np.sum(_compute_log_p_x_i_given_theta(x, theta)).item()
105
106
107 def _compute_log_e_step(x: np.ndarray, theta: Theta) -> np.ndarray:
108     """
109     Compute the e step of expectation maximisation
110     :param x: the image data (n, d)
111     :param theta: the parameters of the model
112     :return: an array of the log responsibilities of k mixture components for each image (n, k)
113     """
114     log_r_unnormalised = _compute_log_component_p_x_i_given_theta(x, theta)
115     log_r_normaliser = logsumexp(log_r_unnormalised, axis=1)
116     log_responsibility = log_r_unnormalised - log_r_normaliser[:, np.newaxis]
117     return log_responsibility
118
119
120 def _compute_log_pi_hat(log_responsibility: np.ndarray) -> np.ndarray:
121     """
122     Compute the log of the maximised mixing proportions
123     :param log_responsibility: an array of the log responsibilities of k mixture components for each image
124     (n, k)
125     :return: an array of the maximised log mixing proportions (1, k)
126     """
127     n, _ = log_responsibility.shape
128     return (logsumexp(log_responsibility, axis=0) - np.log(n)).reshape(1, -1)
129
130
131 def _compute_log_p_matrix_hat(
132     x: np.ndarray, log_responsibility: np.ndarray
133 ) -> np.ndarray:
134     """
135     Compute the log of the maximised pixel probabilities
136     :param x: the image data (n, d)
137     :param log_responsibility: an array of the log responsibilities of k mixture components for each image
138     (n, k)
139     :return: an array of the maximised pixel probabilities for each component (d, k)
140     """
141     n, d = x.shape
142     _, k = log_responsibility.shape
143
144     x_repeated = np.repeat(x[:, :, np.newaxis], k, axis=2) # (n, d, k)
145     log_responsibility_repeated = np.repeat(
146         log_responsibility[:, np.newaxis, :], d, axis=1
147     ) # (n, d, k)
148
149     alpha = 2
150     beta = 2
151
152     log_p_matrix_unnormalised_posterior = logsumexp(
153         log_responsibility_repeated, b=(x_repeated + alpha - 1), axis=0
154     ) # (d, k)
155
156     log_p_matrix_normaliser_posterior = logsumexp(
157         log_responsibility_repeated, b=(alpha + beta - 1), axis=0
158     ) # (d, k)
159
160     log_p_matrix_normalised_posterior = (
161         log_p_matrix_unnormalised_posterior - log_p_matrix_normaliser_posterior
162     )
163     return log_p_matrix_normalised_posterior
164
165
166 def _compute_log_m_step(x: np.ndarray, log_responsibility: np.ndarray) -> Theta:
167     """
168     Compute the m step of expectation maximisation
169     :param x: the image data (n, d)
170     :param log_responsibility: an array of the log responsibilities of k mixture components for each image
171     (n, k)
172     :return: thetas optimised after maximisation step
173     """
174     return Theta(
175         log_pi=_compute_log_pi_hat(log_responsibility),
176         log_p_matrix=_compute_log_p_matrix_hat(x, log_responsibility),
177     )
178
179
180 def _run_expectation_maximisation(
181     x: np.ndarray, theta: Theta, max_number_of_steps: int, epsilon: float
182 ) -> Tuple[Theta, np.ndarray, List[float]]:
183     """
184     Run the expectation maximisation algorithm
185     :param x: the image data (n, d)
186     :param theta: initial theta parameters
187     :param max_number_of_steps: the maximum number of steps to run the algorithm
188     :param epsilon: the minimum required change in log likelihood, otherwise the algorithm stops early
189     :return: a tuple containing the optimised thetas, the log responsibilities,

```

```

187         and the log likelihood at each step of the algorithm
188     """
189     log_responsibility = None
190     log_likelihoods = []
191     for _ in range(max_number_of_steps):
192         log_responsibility = _compute_log_e_step(x, theta)
193         theta = _compute_log_m_step(x, log_responsibility)
194
195         log_likelihoods.append(_compute_log_likelihood(x, theta))
196
197     # check for early stopping
198     if len(log_likelihoods) > 1:
199         if (log_likelihoods[-1] - log_likelihoods[-2]) < epsilon:
200             break
201     return theta, log_responsibility, log_likelihoods
202
203
204 def _plot_p_matrix(
205     thetas: List[Theta], ks: List[int], figure_title: str, figure_path: str
206 ):
207     n = len(ks)
208     m = np.max(ks)
209     fig = plt.figure()
210     fig.set_figwidth(15)
211     fig.set_figheight(10)
212     for i, k in enumerate(ks):
213         for j in range(k):
214             ax = plt.subplot(n, m, m * i + j + 1)
215             ax.imshow(
216                 thetas[i].p_matrix[:, :, j],
217                 interpolation="None",
218             )
219             ax.tick_params(
220                 axis="x",
221                 which="both",
222                 bottom=False,
223                 top=False,
224             )
225             ax.tick_params(
226                 axis="y",
227                 which="both",
228                 left=False,
229                 right=False,
230             )
231             ax.xaxis.set_ticklabels([])
232             ax.yaxis.set_ticklabels([])
233             ax.set_title(f"pi-{j}: {np.round(thetas[i].pi[0, j], 2)}")
234             if j == 0:
235                 ax.set_ylabel(f"{k=}")
236     fig.suptitle(figure_title)
237     plt.savefig(figure_path)
238
239
240 def _plot_tsne_responsibility_clusters(
241     log_responsibilities: List[np.ndarray],
242     ks: List[int],
243     figure_title: str,
244     figure_path: str,
245 ):
246     n = len(ks)
247     fig = plt.figure()
248     fig.set_figwidth(5 * n)
249     fig.set_figheight(5)
250     for i, k in enumerate(ks):
251         if k > 2:
252             embedding = TSNE(
253                 n_components=2,
254                 learning_rate="auto",
255                 init="random",
256                 perplexity=10,
257                 random_state=DEFAULT_SEED,
258             ).fit_transform(log_responsibilities[i])
259         else:
260             embedding = np.exp(log_responsibilities[i])
261         ax = plt.subplot(1, n, i + 1)
262         ax.scatter(embedding[:, 0], embedding[:, 1])
263         ax.set_title(f"{k=}")
264     fig.suptitle(figure_title)
265     plt.savefig(figure_path, bbox_inches="tight")
266
267
268 def _plot_log_posteriors(
269     log_posteriors: List[List[float]],
270     ks: List[int],
271     epsilon: float,
272     figure_title: str,
273     figure_path: str,
274 ) -> None:
275     fig, ax = plt.subplots(len(ks), 1, constrained_layout=True)
276     fig.set_figwidth(10)
277     fig.set_figheight(10)
278     for i, k in enumerate(ks):
279         ax[i].plot(np.arange(1, len(log_posteriors[i]) + 1), log_posteriors[i])
280         ax[i].set_xlabel("Step")
281         ax[i].set_ylabel(f"Log-Posterior")
282         ax[i].set_title(f"{k=} {epsilon=}")

```

```

283 plt.suptitle (figure_title)
284
285 plt.savefig (figure_path)
286
287
288 def e(
289     x: np.ndarray,
290     number_of_trials: int,
291     ks: List[int],
292     epsilon: float,
293     max_number_of_steps: int,
294     figure_path: str,
295     figure_title: str,
296 ) -> None:
297     n, d = x.shape
298     np.random.seed(DEFAULT_SEED)
299     for i in range(number_of_trials):
300         init_thetas = []
301         em_thetas = []
302         log_posteriors = []
303         log_responsibilities = []
304         for j, k in enumerate(ks):
305             init_theta = _init_params(k, d)
306             em_theta, log_responsibility, log_posterior = _run_expectation_maximisation(
307                 x,
308                 theta=init_theta,
309                 epsilon=epsilon,
310                 max_number_of_steps=max_number_of_steps,
311             )
312             init_thetas.append(init_theta)
313             em_thetas.append(em_theta)
314             log_responsibilities.append(log_responsibility)
315             log_posteriors.append(log_posterior)
316
317         _plot_p_matrix(
318             init_thetas,
319             ks,
320             figure_title=f"{figure_title} Trial {i}: Initialised P",
321             figure_path=f"{figure_path}-{i}-initialised-p.png",
322         )
323         _plot_p_matrix(
324             em_thetas,
325             ks,
326             figure_title=f"{figure_title} Trial {i}: EM Optimised P",
327             figure_path=f"{figure_path}-{i}-optimised-p.png",
328         )
329         _plot_tsne_responsibility_clusters(
330             log_responsibilities,
331             ks,
332             figure_title=f"{figure_title} Trial {i}: TSNE Responsibility Visualisation",
333             figure_path=f"{figure_path}-{i}-tsne.png",
334         )
335         _plot_log_posteriors(
336             log_posteriors,
337             ks,
338             epsilon,
339             figure_title=f"{figure_title} Trial {i}: Log-Posterior",
340             figure_path=f"{figure_path}-{i}-log-pos.png",
341         )

```

src/solutions/q3.py

- (e) Running the algorithm a few times starting from randomly chosen initial conditions and visualising the parameters:

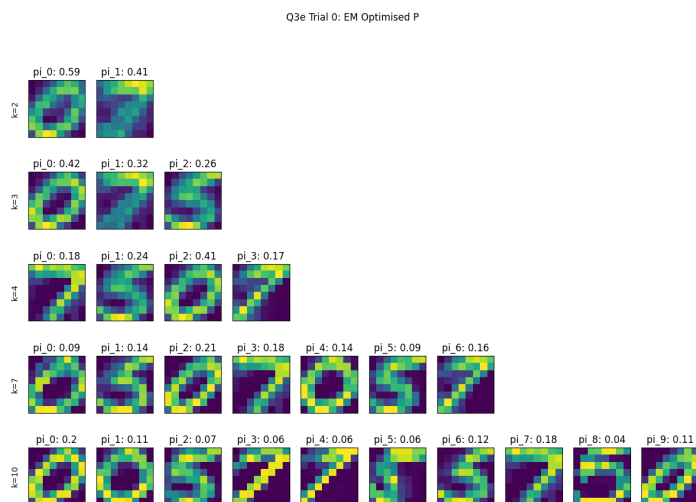


Figure 7: EM optimised parameters: Trial 0

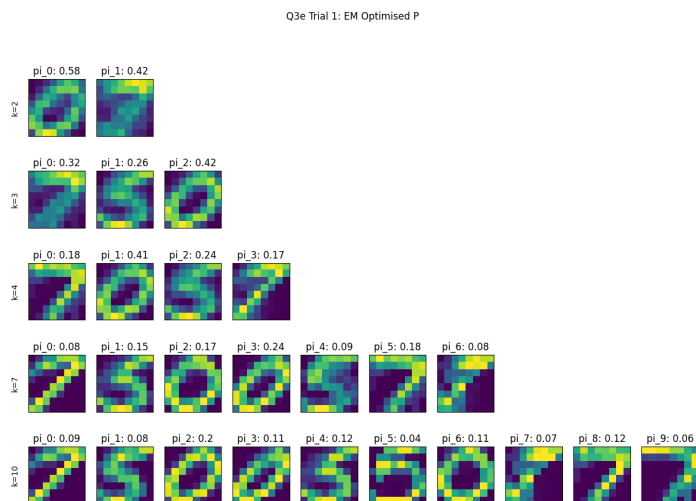


Figure 8: EM optimised parameters: Trial 1

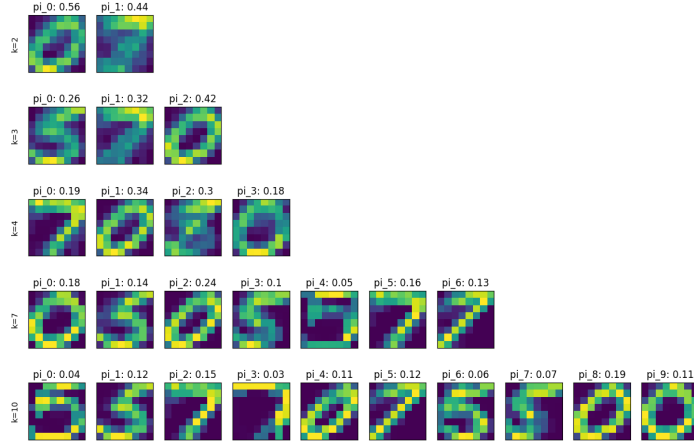


Figure 9: EM optimised parameters: Trial 2

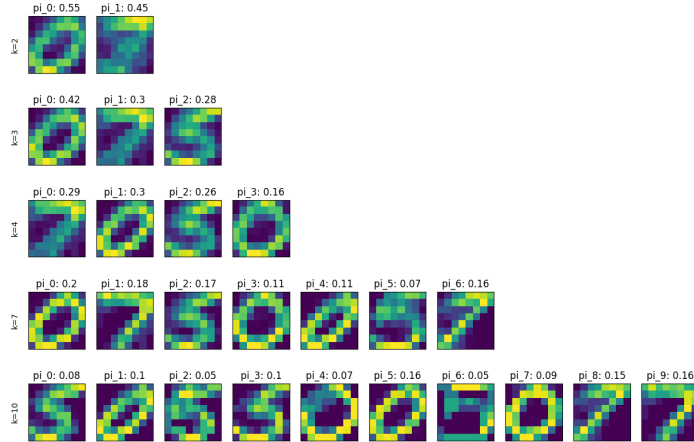


Figure 10: EM optimised parameters: Trial 3

For smaller k , we can visually see that we obtain very similar solutions (a 7 and a 0 for $k = 2$). However for higher K , we see that this may not always be the case. For Trial 2 of $k = 10$, we have three 5's whereas in Trial 4 we have two 5's. Interestingly, different clusters of the same digits can be different, representing different variants of the written digit (i.e. a slanted zero, a slightly slanted zero, and a symmetric zero).

Moreover, looking at the responsibilities of each mixture component, we can see that when k is relatively small they are relatively evenly distributed. However for $k = 7$ and especially $k = 10$, we can see some components have very small or zero probability (i.e. π_2 of trial 2). It will be unlikely for those components to represent very distinct clusters (i.e. the parameters for π_2 and π_9 are very similar in trial 2) This can be verified when we perform a TSNE visualisation of the responsibility vector for each of the images (Note that for $k = 2$, the responsibility vector is displayed). We can see that for large k , qualitatively the number of clusters no longer matches the k value, indicating that some clusters are redundant. For example for $k = 7$ and $k = 10$ we can only qualitatively see four or five clusters with TSNE.

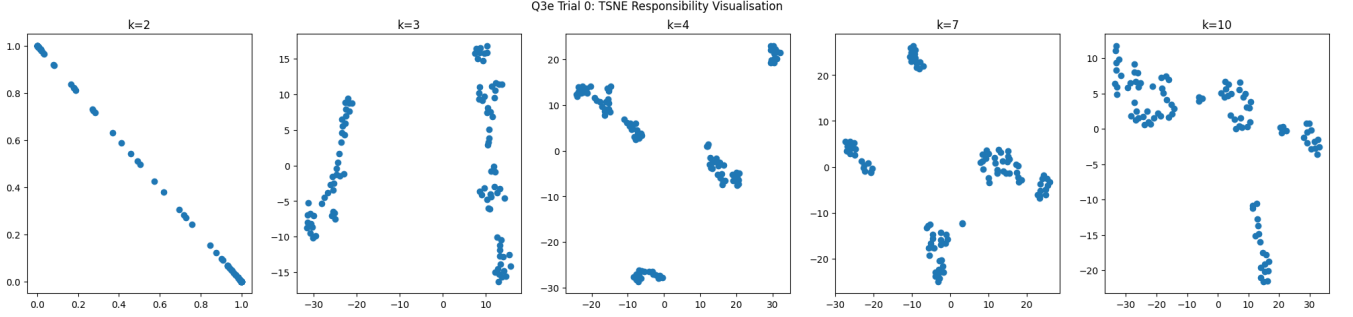


Figure 11: TSNE Visualisation of Image responsibilities: Trial 0

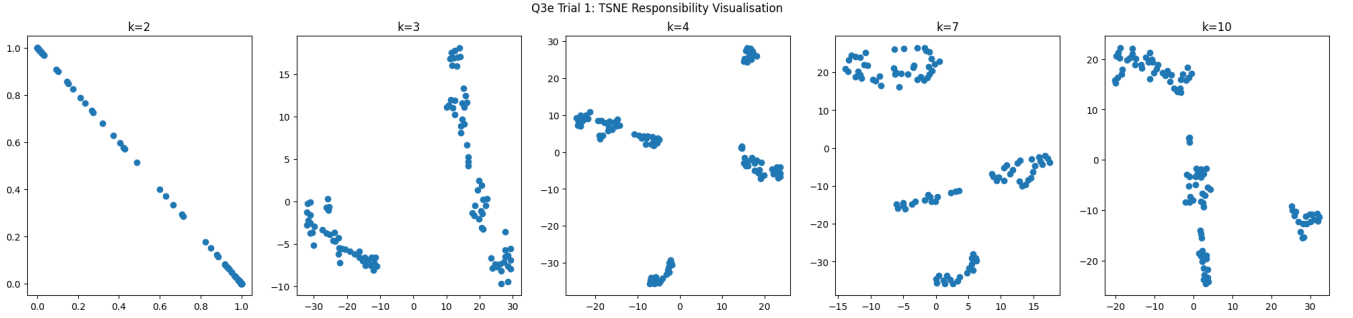


Figure 12: TSNE Visualisation of Image responsibilities: Trial 1

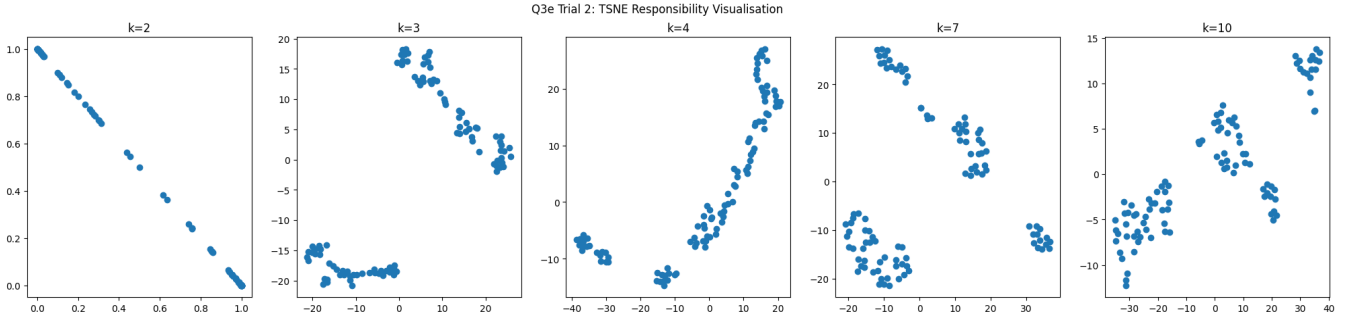


Figure 13: TSNE Visualisation of Image responsibilities: Trial 2

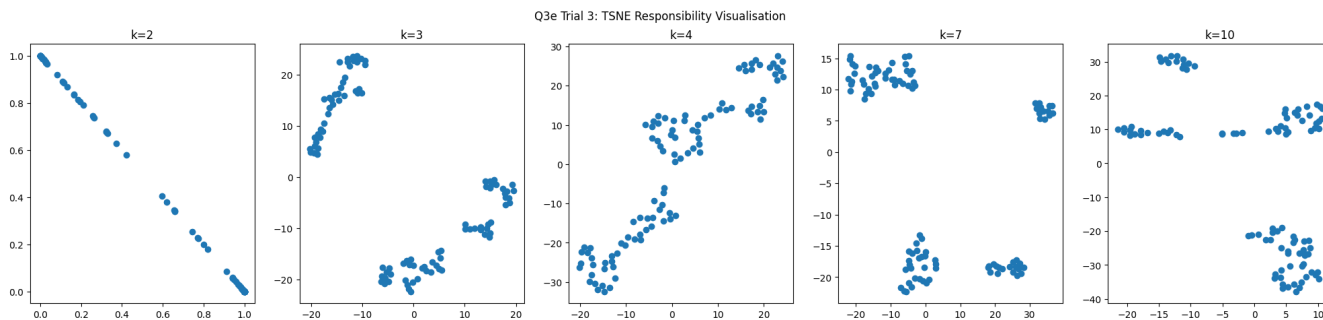


Figure 14: TSNE Visualisation of Image responsibilities: Trial 3

Improvements to the model could include searching for an optimal k by maximising the log posterior with regularisation on the magnitude of k to balancing maximising log posterior with minimising model complexity. Additionally, adding a prior on the responsibility components can be helpful to ensure non-zero mixing components unlike the components visualised here. This could help promote more meaningful clusters as k increases.

[BONUS] Express the log-likelihoods obtained in bits and relate these numbers to the length of the naive encoding of these binary data. How does your number compare to gzip (or another compression algorithm)? Why the difference? [5 marks]

[BONUS] Consider the total cost of encoding both the model parameters and the data given the model. How does this total cost compare to gzip (or similar)? How does it depend on K ? What might this tell you? [5 marks]

Question 5

- (a) The formulae for the ML estimates of $P(s_i = \alpha | s_{i-1} = \beta) = \Psi(\alpha, \beta)$:

$$\Psi(\alpha, \beta) = \frac{N_{s_i, s_{i-1}}}{N_{s_{i-1}}}$$

where $N_{s_i, s_{i-1}}$ is the count of the number of occurrences of the pair (s_i, s_{i-1}) , where s_{i-1} is followed by s_i and $N_{s_{i-1}}$ is the number of occurrences of s_{i-1} .

Moreover, the stationary distribution ϕ can be calculated using the power method:

- (i) Initialise any $\phi_0 \in \mathbb{R}^{53 \times 1}$
- (ii) Repeat $\phi_{i+1} = \Psi \phi_i$
- (iii) Terminate when $\phi_{i+1} - \phi_i < \epsilon$

where $\Psi \in \mathbf{R}^{53 \times 53}$ containing the transition probabilities, $\Psi_{i,j} = P(\alpha_j | \alpha_i)$ where α_i is the i^{th} symbol and α_j is the j^{th} symbol, and ϵ is some small number indicating sufficient convergence of the distribution to be considered stationary. The function $\phi(\gamma)$ is simply the index of γ in the vector ϕ .

The transition matrix Ψ :

[illegible]

(Apologies for the tiny font, latex was being difficult)

The invariant distribution ϕ :

<i>Symbol</i>	<i>Probability</i>
=	1.7e-05
space	1.7e-01
-	6.1e-04
,	1.2e-02
;	3.9e-04
:	2.9e-04
!	6.0e-04
?	4.7e-04
/	1.9e-05
.	7.7e-03
'	1.9e-05
double quotes	2.4e-05
(2.3e-04
)	2.2e-04
[1.7e-05
]	1.7e-05
*	1.1e-04
0	6.9e-05
1	1.4e-04
2	6.0e-05
3	3.4e-05
4	2.3e-05
5	3.2e-05
6	3.2e-05
7	2.8e-05
8	7.6e-05
9	2.6e-05
a	6.6e-02
b	1.1e-02
c	2.0e-02
d	3.8e-02
e	1.0e-01
f	1.8e-02
g	1.6e-02
h	5.4e-02
i	5.6e-02
j	8.5e-04
k	6.4e-03
l	3.1e-02
m	2.0e-02
n	5.9e-02
o	6.2e-02
p	1.5e-02
q	7.7e-04
r	4.7e-02
s	5.2e-02
t	7.2e-02
u	2.1e-02
v	8.5e-03
w	1.9e-02
x	1.4e-03
y	1.5e-02
z	7.4e-04

- (b) The latent variables $\sigma(s)$ for different symbols s are not independent. This is because by choosing an encoding for one symbol $e = \sigma(s)$, the encoding for a second symbol $\sigma(s')$ cannot be e . We have 53 symbols but only 52 degrees of freedom, because once we have defined the encoding for 52 symbols, the encoding for the 53rd symbol cannot be chosen. Thus, there exists a dependence between the symbols for a given σ .

The joint probability of the encrypted text $e_1 e_2 \dots e_n$ given σ :

$$P(e_1, e_2, \dots, e_n | \sigma) = \phi(\gamma = \sigma^{-1}(e_1)) \prod_{i=2}^n \psi(\alpha = \sigma^{-1}(e_i), \beta = \sigma^{-1}(e_{i-1}))$$

because σ is the encoding function, mapping a symbol s into the encoded symbol e , we require σ^{-1} the decoding function mapping the encoded symbol e back to s .

- (c) The proposal probability $S(\sigma \rightarrow \sigma')$ depends on the permutations of σ and σ' . Our proposal generating process restricts us to choose a proposal σ' that differs from σ only at *two* spots:

$$\sigma'(s^i) = \sigma(s^j)$$

$$\sigma'(s^j) = \sigma(s^i)$$

for any two symbols s^i and s^j of the 53 possible symbols ($s^i \neq s^j$).

Therefore, if the above doesn't hold for σ' , $S(\sigma \rightarrow \sigma') = 0$. From σ there are $\binom{53}{2}$ possible proposal σ' 's with the above property. Because we are assuming a uniform prior distribution over σ 's, the transition probability of a σ' that satisfies the above property is $S(\sigma \rightarrow \sigma') = \frac{1}{\binom{53}{2}}$.

The MH acceptance probability is given as:

$$A(\sigma \rightarrow \sigma' | \mathcal{D}) = \min\left\{1, \frac{S(\sigma' \rightarrow \sigma)P(\sigma' | \mathcal{D})}{S(\sigma \rightarrow \sigma')P(\sigma | \mathcal{D})}\right\}$$

because $S(\sigma \rightarrow \sigma')$ is the conditional transition probability of σ' given σ and \mathcal{D} is our encrypted text e_1, e_2, \dots, e_n .

$S(\sigma \rightarrow \sigma') = S(\sigma' \rightarrow \sigma)$ for all σ and σ' that differ only at two spots because the probability in this case will always be $\frac{1}{\binom{53}{2}}$, we can simplify:

$$A(\sigma \rightarrow \sigma' | \mathcal{D}) = \min\left\{1, \frac{P(\sigma' | \mathcal{D})}{P(\sigma | \mathcal{D})}\right\}$$

From Bayes' Theorem:

$$P(\sigma | \mathcal{D}) = \frac{P(\mathcal{D} | \sigma)P(\sigma)}{\sum_{\sigma'} P(\mathcal{D} | \sigma')P(\sigma')}$$

We are assuming a uniform prior for σ , so $P(\sigma)$ is a constant and we can simplify further:

$$A(\sigma \rightarrow \sigma' | \mathcal{D}) = \min\left\{1, \frac{P(\mathcal{D} | \sigma')}{P(\mathcal{D} | \sigma)}\right\}$$

This is the acceptance probability for a given proposal σ' . The expression for $P(\mathcal{D} | \sigma)$ is $P(e_1, e_2, \dots, e_n | \sigma)$ described in the previous part.

(d) Reporting the current decryption of the first 60 symbols after every 100 iterations:

MH Iteration	Current Decryption
0	d0[?0?sdhrg0tde0]gr0as[drgti]r0?rtg 0[?0bt4org0ltar0]r0[? ,
100	odzhvzv.sdfirtz d5zh,trzasldrt ilrvzr twzhvzb egtzf arzhzw,
200	odgh.g.sdfirtz dagh,trg/sldrt blrg.r tuhg.gi e?rtgf /rghgu,
300	odrh.r.sdfirtz darh-igrksldgi blgr.g iurh.rt eygirt kgrhgru-
400	idrkrlrs-sdfior dark-ogr.sldgo blgrhg ourkhrt eygorf .grkgru-
500	idrkrlrs-sdfior dark-olnsgsldo bgrlhl ourkhrt eylorf ulrkdu-
600	idrwlrhasdlofr d-rwaforpsgdof bgorho furwhrt eyofrl porworua
700	idrwlrhasdlofr d-rwaforpsgdof bgorho turwhrf eyotrl porworua
800	idrwlrhasd-otr dlrwaforpsgdof bgorho turwhrf eyotr- porworua
900	idrwlrhasdgotr dlrwaforpsgdof bcorho tprwhrf eyotrg uorworpa
1000	idrwlrhasd-otr dlrwaforpsgdof bcorho tprwhrf eyotr- gorworpa
1100	ilrwlrhasl.otr ldrwaforpsgdof ncorho tfrwhrf eyotr- gorworfa
1200	ilrwlrhasl.otr ldrwaforpsgdof ncorho tfrwhrf eyotr- gorworfa
1300	ilrwlrhasl.ofr ldrwaforpsgdof ncorho tfrwhrf eyofr- gorworta
1400	ilrwlrhasl.ofr ldrwaforpsgdof ncorho tfrwhrf eyofr- gorworta
1500	inrwlrhasngofr ndrwafor.scnof bcorho tfrwhrf eyofrg .orworta
1600	inrchrhasngofr ndrcafor.swnof bcorho tfrchrf eyofrg .orcorta
1700	inrchrhasngofr ndrcafor.swnof bcorho tfrchrf eyofrg .orcorta
1800	inrchrhasngofr ndrcafor.swnof bcorho tfrchrf eyofrg lororta
1900	inrchrhasngofr ndrcafor.swnof bcorho tfrchrf eyofrg lororta
2000	inrchrhasngofr ndrcafor.swnof bcorho tfrchrf eyofrg lororta
2100	inechehasngofe ndecafoelswnof bworho ftrchep ryofeg loecoeta
2200	in ch hasn.of end calo lsnolevwo hoelt ch geryof .elo co ta
2300	in wh hasn.of end wafo lsnolevwo hoelt wh geryof .elo wo ta
2400	in wh hasn.of end wafo lsnolevwo hoelt wh geryof .elo wo ta
2500	in wh hasn.of end wafo lsnolevwo hoelt wh geryof .elo wo ta
2600	in wh hasn.of end wafo lsnolevwo hoelt wh geryof .elo wo ta
2700	in wh hasn.of end wafo lsnolevwo hoelt wh geryof .elo wo ta
2800	in ch hasn.ol end calo fswnolevwo hoelt ch geryof .elo co ta
2900	in ch hasn.ol end calo fswnolevwo hoelt ch geryof .elo co ta
3000	in ch hasn.ol end calo fswnolevwo hoelt ch geryof .elo co ta
3100	in ch haun.ol end calo fuwnolevwo hoelt ch geryof .elo co ta
3200	in ch haun.ol end calo fuwnolevwo hoelt ch geryof .elo co ta
3300	in ch haun.os end caso fuwnosevwo hoelt ch geryos .elo co ta
3400	in ch haun.os end caso fuwnosevwo hoelt ch geryos .elo co ta
3500	in ch haun.os end caso fuwnosevwo hoelt ch geryos .elo co ta
3600	in cy yaun.er end caro fuwnorevwo yeort cy geshor .elo co ta
3700	in cy yaun.er end care lubnerovwe yeort cy gosher .ole ce ta
3800	in cy yaun.er end care lubnerovwe yeort cy fosher .ole ce ta
3900	in cy yaun.er end care lubnerovwe yeort cy fosher .ole ce ta
4000	in cy yaun.er end care bufnervof yeort cy losher .obe ce ta
4100	in my yaun.er end mare bufnervof yeort my losher .obe me ta
4200	in my yaun.er end mare bufnervof yeort my losher .obe me ta
4300	in my yaun.er end mare bufnervof yeort my losher .obe me ta
4400	in my yaun.er end mare bufnervof yeort my losher .obe me ta
4500	in my yaun.er end mare vufnerobfe yeort my losher .ove me ta
4600	in my yaun.er end mare vufnerobfe yeort my losher .ove me ta
4700	in my yaun.er end mare vufnerobfe yeort my losher .ove me ta
4800	in my yaunger end mare vufnerobfe yeors my lother gove me sa
4900	in my yaunger end mare vufnerobfe yeors my lother gove me sa
5000	in my yaunger end mare vufnerobfe yeors my lother gove me sa
5100	in my yaunger end mare vufnerobfe yeors my lother gove me sa
5200	in my yaunger end mare vufnerobfe yeors my lother gove me sa
5300	in my yaunger end mare vufnerobfe yeors my lother gove me sa
5400	in my yaunger end mare vufnerobfe yeors my lother gove me sa
5500	in my yaunger end mare vufnerobfe yeors my lother gove me sa
5600	in my yaunger end mare vufnerobfe yeors my lother gove me sa
5700	in my yaunger end mare vufnerobfe yeors my lother gove me sa
5800	in my yaunger end mare vufnerobfe yeors my lother gove me sa
5900	in my yaunger end mare vufnerobfe yeors my lother gove me sa
6000	in my yaunger end mare vufnerobfe yeors my lother gove me sa
6100	in my yaunger end mare vufnerobfe yeors my lother gove me sa
6200	in my yaunger end mare vufnerobfe yeors my lother gove me sa
6300	in my yaunger end mare vufnerobfe yeors my lother gove me sa
6400	in my yaunger end mare vufnerobfe yeors my fother gove me sa
6500	in my yaunger end mare vufnerobfe yeors my fother gove me sa
6600	in my younger end more vulnerable yeors my father gave me so
6700	in my younger end more vulnerable yeors my father gave me so
6800	in my younger end more vulnerable yeors my father gave me so
6900	in my younger end more vulnerable yeors my father gave me so
7000	in my younger end more vulnerable yeors my father gave me so
7100	in my younger end more vulnerable yeors my father gave me so
7200	in my younger end more vulnerable yeors my father gave me so
7300	in my younger end more vulnerable yeors my father gave me so
7400	in my younger end more vulnerable yeors my father gave me so
7500	in my younger end more vulnerable yeors my father gave me so
7600	in my younger end more vulnerable yeors my father gave me so
7700	in my younger end more vulnerable yeors my father gave me so
7800	in my younger end more vulnerable yeors my father gave me so
7900	in my younger end more vulnerable yeors my father gave me so
8000	in my younger end more vulnerable yeors my father gave me so
8100	in my younger end more vulnerable yeors my father gave me so
8200	in my younger end more vulnerable yeors my father gave me so
8300	in my younger end more vulnerable yeors my father gave me so
8400	in my younger end more vulnerable yeors my father gave me so
8500	in my younger end more vulnerable yeors my father gave me so
8600	in my younger end more vulnerable yeors my father gave me so
8700	in my younger end more vulnerable yeors my father gave me so
8800	in my younger end more vulnerable yeors my father gave me so
8900	in my younger end more vulnerable yeors my father gave me so
9000	in my younger end more vulnerable yeors my father gave me so
9100	in my younger end more vulnerable yeors my father gave me so
9200	in my younger end more vulnerable yeors my father gave me so
9300	in my younger end more vulnerable yeors my father gave me so
9400	in my younger end more vulnerable yeors my father gave me so
9500	in my younger end more vulnerable yeors my father gave me so
9600	in my younger end more vulnerable yeors my father gave me so
9700	in my younger end more vulnerable yeors my father gave me so
9800	in my younger end more vulnerable yeors my father gave me so
9900	in my younger end more vulnerable yeors my father gave me so
10000	in my younger end more vulnerable yeors my father gave me so

The corresponding σ :

s	$\sigma(s)$
=	(
space	x
-	h
,	,
;	l
:	n
!	r
?	e
/	f
.	b
'	2
double quotes	double quotes
(3
)	=
[i
]	o
*	l
0	z
1	m
2	c
3	8
4)
5	.
6	*
7	k
8	0
9	q
a	/
b	:
c	-
d	;
e	5
f	6
g	s
h	9
i	'
j]
k	[
l	y
m	v
n	d
o	4
p	space
q	?
r	g
s	t
t	7
u	p
v	j
w	a
x	u
y	!
z	w

To help with chain initialisation, 10000 different σ 's were randomly and independently sampled. The σ providing the best log-likelihood was chosen as the starting point for the MH chain and algorithm was then run for 10000 iterations. Moreover, ten different trials were performed, where the trial with the best log-likelihood is displayed.

The Python code for the MH sampler:

```
1 from typing import Dict, List, Tuple
2
3 import numpy as np
4 import pandas as pd
5 from sklearn.preprocessing import normalize
6
7 from src.constants import DEFAULT_SEED
8
9
10 class Decrypter:
11     def __init__(self, decryption_dict):
12         self.decryption_dict = decryption_dict
13
14     def decrypt(self, encrypted_message):
15         return "".join([self.decryption_dict[x] for x in encrypted_message])
16
17
18 class Statistics:
19     def __init__(
20         self,
21         training_text: str,
22         symbols: List[str],
23         invariant_stopping_epsilon: float = 5e-20,
24     ):
25         self.training_text = training_text
26         self.symbols = symbols
27         self.num_symbols = len(symbols)
28         self.symbols_dict = {k: v for v, k in enumerate(symbols)}
29         self.text_numbers = [
30             self.symbols_dict[symbol]
31             for symbol in list(training_text)
32             if symbol in self.symbols_dict
33         ]
34         self.transition_matrix = self._construct_transition_matrix(
35             training_text, self.symbols_dict
36         )
37         self.invariant_distribution = self._approximate_invariant_distribution(
38             invariant_stopping_epsilon
39         )
40         self.log_transition_matrix = np.log(self.transition_matrix)
41         self.log_invariant_distribution = np.log(self.invariant_distribution)
42
43     def _construct_transition_matrix(
44         self, training_text: str, symbols_dict: Dict[str, int]
45     ) -> np.ndarray:
46
47         # initialise with ones to ensure ergodicity
48         transition_matrix = np.ones((self.num_symbols, self.num_symbols))
49         for i in range(1, len(training_text)):
50             # check symbols are valid
51             if (
52                 training_text[i] in symbols_dict
53                 and training_text[i - 1] in symbols_dict
54             ):
55                 transition_matrix[
56                     symbols_dict[training_text[i - 1]], symbols_dict[training_text[i]]
57                 ] += 1
58         # normalise to get transition probabilities
59         transition_matrix = normalize(transition_matrix, axis=0, norm="l1")
60         return transition_matrix
61
62     def _approximate_invariant_distribution(
63         self, invariant_stopping_epsilon: float
64     ) -> np.ndarray:
65         invariant_distribution = np.zeros((self.num_symbols, 1))
66         previous_invariant_distribution = invariant_distribution.copy()
67         invariant_distribution[0] = 1
68
69         while (
70             np.linalg.norm(invariant_distribution - previous_invariant_distribution)
71             > invariant_stopping_epsilon
72         ):
73             previous_invariant_distribution = invariant_distribution.copy()
74             invariant_distribution = self.transition_matrix @ invariant_distribution
75         return invariant_distribution
76
77     def log_transition_probability(self, alpha: str, beta: str) -> float:
78         return self.log_transition_matrix[
79             self.symbols_dict[beta], self.symbols_dict[alpha]
80         ]
81
82     def log_invariant_probability(self, gamma: str) -> float:
83         return self.log_invariant_distribution[self.symbols_dict[gamma]].item()
84
85     def compute_log_probability(self, message: str) -> float:
86         log_probability = self.log_invariant_probability(message[0])
87         for i in range(1, len(message)):
88             s_i = message[i]
89             s_i_minus_1 = message[i - 1]
90             log_probability += self.log_transition_probability(s_i, s_i_minus_1)
91         return log_probability
92
93
94 class MetropolisHastingsDecryption:
```

```

95 def __init__(self, symbols):
96     self.symbols = symbols
97     self._random_generator = np.random.default_rng()
98
99 def generate_random_decrypter(self) -> Decrypter:
100     return Decrypter(
101         {
102             self.symbols[i]: self.symbols[x]
103             for i, x in enumerate(
104                 np.random.permutation(np.arange(len(self.symbols)))
105             )
106         }
107     )
108
109 @staticmethod
110 def generate_proposal_decrypter(decrypter: Decrypter) -> Decrypter:
111     x1 = np.random.choice(list(decrypter.decryption_dict.keys()))
112     x2 = np.random.choice(list(decrypter.decryption_dict.keys()))
113     proposal_decryption = decrypter.decryption_dict.copy()
114     proposal_decryption[x2], proposal_decryption[x1] = (
115         decrypter.decryption_dict[x1],
116         decrypter.decryption_dict[x2],
117     )
118     return Decrypter(proposal_decryption)
119
120 def _choose_decrypter(
121     self,
122     statistics,
123     encrypted_message,
124     current_decrypter: Decrypter,
125     proposal_decrypter: Decrypter,
126 ) -> Decrypter:
127     current_log_probability = statistics.compute_log_probability(
128         message=current_decrypter.decrypt(encrypted_message),
129     )
130     proposal_log_probability = statistics.compute_log_probability(
131         message=proposal_decrypter.decrypt(encrypted_message),
132     )
133     acceptance_probability = np.min(
134         [1, np.exp(proposal_log_probability - current_log_probability)]
135     )
136     return self._random_generator.choice(
137         [current_decrypter, proposal_decrypter],
138         p=[1 - acceptance_probability, acceptance_probability],
139     )
140
141 def _find_good_starting_decrypter(
142     self,
143     statistics: Statistics,
144     encrypted_message,
145     number_start_attempts,
146 ) -> Decrypter:
147     best_log_likelihood = -np.float("inf")
148     best_decrypter = None
149     for _ in range(number_start_attempts):
150         decrypter = self.generate_random_decrypter()
151         if (
152             statistics.compute_log_probability(
153                 message=decrypter.decrypt(encrypted_message)
154             )
155             > best_log_likelihood
156         ):
157             best_decrypter = decrypter
158     return best_decrypter
159
160 def run(
161     self,
162     encrypted_message: str,
163     statistics: Statistics,
164     number_of_mh_loops: int,
165     number_start_attempts: int,
166     check_decryption_interval: int,
167     check_decryption_size: int,
168 ) -> Tuple[Decrypter, List[str]]:
169     decrypter = self._find_good_starting_decrypter(
170         statistics, encrypted_message, number_start_attempts
171     )
172     logged_decryption_message = [
173         decrypter.decrypt(encrypted_message)[:check_decryption_size]
174     ]
175     for i in range(1, number_of_mh_loops + 1):
176         if (i + 1) % check_decryption_interval == 0:
177             logged_decryption_message.append(
178                 decrypter.decrypt(encrypted_message)[:check_decryption_size]
179             )
180             proposal_decrypter = self.generate_proposal_decrypter(decrypter)
181             decrypter = self._choose_decrypter(
182                 statistics, encrypted_message, decrypter, proposal_decrypter
183             )
184     return decrypter, logged_decryption_message
185
186
187 def _convert_to_scientific_notation(x: float) -> str:
188     return "{:.1e}".format(float(x))
189
190

```



```

191 def a(
192     symbols: List[str],
193     training_text: str,
194     transition_matrix_path: str,
195     invariant_distribution_path: str,
196 ):
197     statistics = Statistics(
198         training_text,
199         symbols,
200     )
201     symbols_for_df = statistics.symbols.copy()
202     symbols_for_df[symbols_for_df.index(" ")] = "space"
203     symbols_for_df[symbols_for_df.index("'")] = "double quotes"
204     df = pd.DataFrame(
205         data=statistics.transition_matrix,
206         columns=symbols_for_df,
207     )
208     df.index = symbols_for_df
209     df.applymap(_convert_to_scientific_notation).to_csv(transition_matrix_path)
210
211     df = (
212         pd.DataFrame(
213             data=statistics.invariant_distribution.reshape(1, -1),
214             columns=symbols_for_df,
215         )
216         .applymap(_convert_to_scientific_notation)
217         .transpose()
218         .reset_index()
219     )
220     df.columns = ["Symbol", "Probability"]
221     df.set_index("Symbol").to_csv(invariant_distribution_path, sep="|")
222
223
224 def d(
225     encrypted_message: str,
226     symbols: List[str],
227     training_text: str,
228     number_trials: int,
229     number_of_mh_loops: int,
230     number_start_attempts: int,
231     check_decryption_interval: int,
232     check_decryption_size: int,
233     decryptor_table_path: str,
234     decrypted_message_iterations_table_path: str,
235 ):
236     statistics = Statistics(
237         training_text,
238         symbols,
239     )
240     np.random.seed(DEFAULT_SEED)
241     metropolis_hastings_decryption = MetropolisHastingsDecryption(symbols)
242     decrypters = []
243     log_likelihoods = []
244     logged_decryption_messages = []
245     decryption_messages = []
246     for i in range(number_trials):
247         (decrypter, logged_decryption_message) = metropolis_hastings_decryption.run(
248             encrypted_message,
249             statistics,
250             number_of_mh_loops,
251             number_start_attempts,
252             check_decryption_interval,
253             check_decryption_size,
254         )
255         decrypters.append(decrypter)
256         log_likelihoods.append(
257             statistics.compute_log_probability(
258                 decrypter.decrypt(encrypted_message)
259             )
260         )
261         logged_decryption_messages.append(logged_decryption_message)
262         decryption_messages.append(
263             decrypter.decrypt(encrypted_message)[:check_decryption_size]
264         )
265
266     # sort trials by log likelihood
267     best_trial = np.argmax(log_likelihoods)
268
269     decrypter_table = pd.DataFrame(
270         decrypters[best_trial].decryption_dict.items(), columns=["s", "sigma(s)"]
271     )
272     decrypter_table[decrypter_table == " "] = "space"
273     decrypter_table[decrypter_table == "'"] = "double quotes"
274     decrypter_table.set_index("s").to_csv(decryptor_table_path, sep="|")
275
276     decrypted_message_iterations_table = pd.DataFrame(
277         [
278             np.arange(0, len(logged_decryption_messages[best_trial]))
279             * check_decryption_interval,
280             logged_decryption_messages[best_trial],
281         ]
282     ).transpose()
283     decrypted_message_iterations_table.columns = ["MH Iteration", "Current Decryption"]
284     decrypted_message_iterations_table.set_index("MH Iteration").to_csv(
285         decrypted_message_iterations_table_path, sep="|")
286

```

src/solutions/q5.py

- (e) When some values of $\Psi(\alpha, \beta) = 0$, this affects the ergodicity of the chain. An ergodic chain is one that is irreducible (i.e. all possible transitions between symbols have probability greater than zero). If $\Psi(\alpha, \beta) = 0$, this means that there is zero probability that β will transition to α , breaking our definition. To restore ergodicity, we can add a small transition probability between all symbols of the chain. This essentially acts as a prior, stating that the probability of a symbol to transition to any other symbol (including itself) should never be zero.
- (f) Analyse this approach to decoding. For instance, would symbol probabilities alone (rather than transitions) be sufficient? If we used a second order Markov chain for English text, what problems might we encounter? Will it work if the encryption scheme allows two symbols to be mapped to the same encrypted value? Would it work for Chinese with > 10000 symbols? [13 marks]

If we were to use symbol probabilities alone for decoding, the joint probability would be:

$$P(e_1, e_2, \dots, e_n | \sigma) = \prod_{i=1}^n P(\sigma^{-1}(e_i))$$

the product of the likelihoods of the decoded letters. In this case, the optimal decoding would simply replace the most frequent symbols in the encrypted message with the most frequent symbols in the training text. This is much more difficult because each letter is assumed to be independent of its neighbours. For a first order Markov chain, we exploit the structure of language by considering pairs of letters. Assuming that as the training text size approaches infinity and the size of the encrypted message also approaches infinity, that the two will have the same symbol frequency and that the probability of each symbol is unique, (i.e. two different decodings can't have the same likelihood), then using symbol probabilities alone should theoretically work. However, in practise we would unlikely to be able to make these assumptions about symbol frequencies from the size of our training set and encrypted message.

A second-order chain should also work in theory. However, this approach is probably practically more difficult for finding a suitable decoding. This is because our transition matrix would contain N^3 , where N is the number of symbols, to account for all possible second order transitions. Our training text would need to increase quadratically to maintain the same ratio of possible transitions to example transitions (number of second order transitions in a text of length N is $N - 2$ and third order its $N - 3$).

For an encryption scheme where two symbols map to the same encrypted value:

$$\exists \alpha, \beta, \sigma(\alpha) = \sigma(\beta), \alpha \neq \beta$$

this approach can become much more complicated. Our $\sigma^{-1}(e)$ is ill-defined, and therefore how we computing the joint probability of the encrypted text is no longer immediately clear. Moreover, generating proposal encodings is not as simple as swapping the encryption for two symbols. This is because we do not know which two symbols map to the same encrypted symbol and simply swapping would preserve the same collision mapping of the current encoding. Overall, many changes would need to be made to the approach to accommodate for these complications. It is not immediately obvious how current approach could work for this case.

If we used this approach for Chinese with ≥ 10000 symbols, we would be attempting to solve the same problem but with $N \geq 10000$ instead of $N = 53$. Similar to the second order Markov chain, although this is theoretically possible, it would require a transition matrix of size $\geq 10000^2$ which is quite impractical. An alternative set up could be with using Chinese phonetics, for which there are likely much fewer than 10000, however this would require a mapping from a phonetic to an encrypted phonetic.

Question 7

- (a) To find the local extrema of the function $f(x, y) = x + 2y$ subject to the constraint $y^2 + xy = 1$, first we define $g(x, y)$:

$$g(x, y) = y^2 + xy - 1$$

where $g(x, y) = 0$ is an equivalent representation of the given constraint.

We can therefore construct the optimisation problem:

$$\min_{\mathbf{x}} f(\mathbf{x})$$

such that $g(\mathbf{x}) = 0$ and $\mathbf{x} := [x, y]^T$.

We can calculate $\nabla f(\mathbf{x})$:

$$\nabla f(\mathbf{x}) = \left[\frac{\partial}{\partial x}(x + 2y), \frac{\partial}{\partial y}(x + 2y) \right]^T$$

$$\nabla f(\mathbf{x}) = [1, 2]^T$$

and calculating $\nabla g(\mathbf{x})$:

$$\nabla g(\mathbf{x}) = \left[\frac{\partial}{\partial x}(y^2 + xy - 1), \frac{\partial}{\partial y}(y^2 + xy - 1) \right]^T$$

$$\nabla g(\mathbf{x}) = [y, 2y + x]^T$$

Solving the constraint optimisation problem with Lagrange multipliers, we set up the equations:

$$\nabla f(\mathbf{x}) + \lambda \nabla g(\mathbf{x}) = \mathbf{0}$$

and

$$g(\mathbf{x}) = 0$$

Giving us the three equations:

$$1 + \lambda y = 0$$

$$2 + \lambda(2y + x) = 0$$

$$y^2 + xy - 1 = 0$$

Substituting $y = \frac{-1}{\lambda}$ from the first equation into the second equation:

$$2 + \frac{-1}{\lambda}(2y + x) = 0$$

$$\frac{-x}{y} = 0$$

We see that $x = 0$. Solving for y in our third equation with $x = 0$:

$$y^2 - 1 = 0$$

We see that $y = \pm 1$ and from the first equation $\lambda \mp 1$.

The local extrema are $(x = 0, y = 1)$ when our $\lambda = -1$ and $(x = 0, y = -1)$ when our $\lambda = 1$.

(b)

(i) Given that $g(a) = \ln(a)$, we want to transform this to the form $f(x, a) = 0$:

$$x = \ln(a)$$

$$\exp(x) - a = 0$$

Thus,

$$f(x, a) = \exp(x) - a$$

(ii) We know that for Newton's method's

$$x_{n+1} = x_n - \frac{f(x_n)}{f'(x_n)}$$

where $f(x_n) = \exp(x_n) - a$

We can calculate:

$$f'(x) = \frac{\partial f(x, a)}{\partial x} = \exp(x)$$

Assuming we can evaluate $\exp(x)$, our update equation:

$$x_{n+1} = x_n - \frac{\exp(x_n) - a}{\exp(x_n)}$$

Simplifying:

$$x_{n+1} = x_n + \frac{a}{\exp(x_n)} - 1$$

Appendix: main.py

```
1 import os
2
3 import numpy as np
4
5 from src.constants import (
6     BINARY_DIGITS_FILE_PATH,
7     MESSAGE_FILE_PATH,
8     OUTPUTS_FOLDER,
9     SYMBOLS_FILE_PATH,
10    TRAINING_TEXT_FILE_PATH,
11 )
12 from src.solutions import q1, q2, q3, q5
13
14 if __name__ == "__main__":
15
16     if not os.path.exists(OUTPUTS_FOLDER):
17         os.makedirs(OUTPUTS_FOLDER)
18
19     x = np.loadtxt(BINARY_DIGITS_FILE_PATH)
20     # Question 1
21     Q1.OUTPUT_FOLDER = os.path.join(OUTPUTS_FOLDER, "q1")
22     if not os.path.exists(Q1.OUTPUT_FOLDER):
23         os.makedirs(Q1.OUTPUT_FOLDER)
24
25     q1.d(
26         x,
27         figure_path=os.path.join(Q1.OUTPUT_FOLDER, "q1d.png"),
28         figure_title="Q1d: Maximum Likelihood Estimate",
29     )
30     q1.e(
31         x,
32         alpha=3,
33         beta=3,
34         figure_path=os.path.join(Q1.OUTPUT_FOLDER, "q1e"),
35         figure_title="Q1e: Maximum A Prior",
36     )
37
38     # Question 2
39     Q2.OUTPUT_FOLDER = os.path.join(OUTPUTS_FOLDER, "q2")
40     if not os.path.exists(Q2.OUTPUT_FOLDER):
41         os.makedirs(Q2.OUTPUT_FOLDER)
42     q2.c(x, table_path=os.path.join(Q2.OUTPUT_FOLDER, "q2c.csv"))
43
44     # Question 3
45     Q3.OUTPUT_FOLDER = os.path.join(OUTPUTS_FOLDER, "q3")
46     if not os.path.exists(Q3.OUTPUT_FOLDER):
47         os.makedirs(Q3.OUTPUT_FOLDER)
48     q3.e(
49         x,
50         number_of_trials=4,
51         ks=[2, 3, 4, 7, 10],
52         epsilon=1e-1,
53         max_number_of_steps=int(1e2),
54         figure_path=os.path.join(Q3.OUTPUT_FOLDER, "q3e"),
55         figure_title="Q3e",
56     )
57
58     # Question 5
59     Q5.OUTPUT_FOLDER = os.path.join(OUTPUTS_FOLDER, "q5")
60     if not os.path.exists(Q5.OUTPUT_FOLDER):
61         os.makedirs(Q5.OUTPUT_FOLDER)
62
63     with open(TRAINING_TEXT_FILE_PATH) as fp:
64         training_text = fp.read().replace("\n", "").lower()
65     with open(SYMBOLS_FILE_PATH) as fp:
66         symbols = fp.read().split("\n")
67     with open(MESSAGE_FILE_PATH) as fp:
68         encrypted_message = fp.read()
69
70     q5.a(
71         symbols,
72         training_text,
73         transition_matrix_path=os.path.join(Q5.OUTPUT_FOLDER, "q5a-transition.csv"),
74         invariant_distribution_path=os.path.join(Q5.OUTPUT_FOLDER, "q5a-invariant.csv"),
75     )
76
77     q5.d(
78         encrypted_message,
79         symbols,
80         training_text,
81         number_trials=10,
82         number_of_mh_loops=int(1e4),
83         number_start_attempts=int(1e4),
84         check_decryption_interval=100,
85         check_decryption_size=60,
86         decryptor_table_path=os.path.join(Q5.OUTPUT_FOLDER, "q5d-decrypter.csv"),
87         decrypted_message_iterations_table_path=os.path.join(
88             Q5.OUTPUT_FOLDER, "q5d-iterations.csv"
89         ),
90     )
```

main.py